



AWS
re:Invent

B L C 2 0 9

Asset provenance ledger system based on Amazon QLDB: BMW's use case

Michael Labib

Principal SA, QLDB
Amazon Web Services

Emil Djerekarov

Technical Architect
BMW Group

Agenda

- Blockchain and purpose-built databases at AWS
- Amazon QLDB architecture and features
- BMW: Digital passport

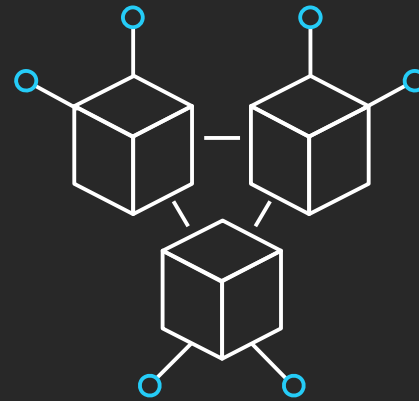
What is blockchain

Ledgers



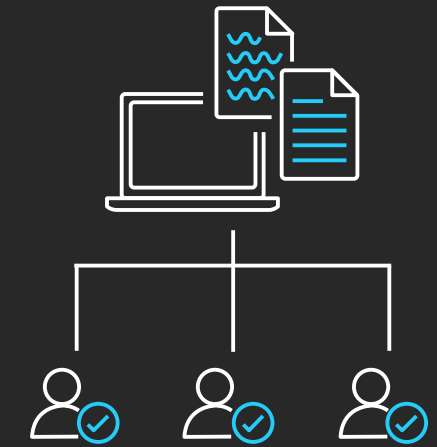
Immutable, append-only,
cryptographically verifiable

Decentralization



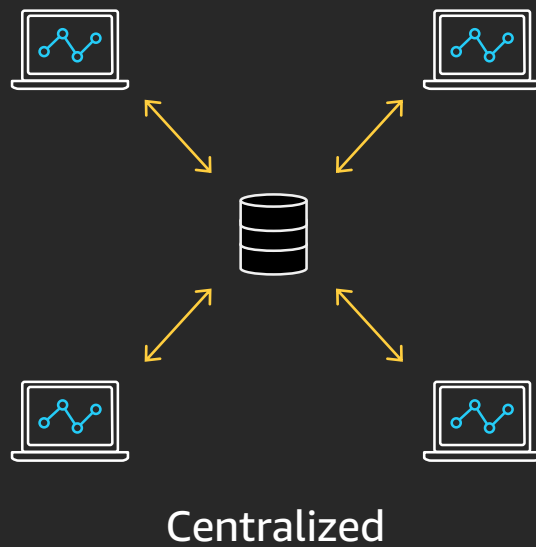
Distributed trust and
data replication

Consensus algorithms

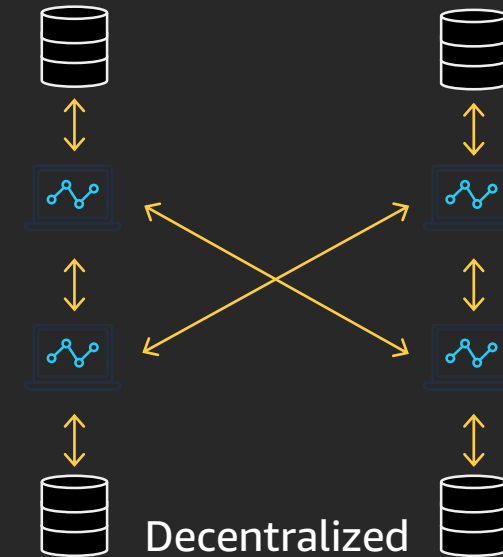


No intermediaries in
decision process, support
for smart contracts

Centralized versus decentralized



- Owned by a single, trusted authority
- Addresses core need of an immutable and verifiable transactional log
- Fast—doesn't require consent from members to commit transactions



- No single owner of the ledger; joint ownership by multiple parties
- Addresses core need of enabling multiple parties to interact transparently with trust
- Removes intermediaries when a group of members needs to transact; can make business processes more efficient

Purpose-built databases at AWS



Relational

Referential integrity,
ACID transactions,
schema-on-write

Lift and shift, ERP,
CRM, finance



Key value

High throughput,
low-latency
reads and writes,
endless scale

Real-time bidding,
shopping cart, social,
product catalog,
customer preferences



Document

Store documents
and quickly access
querying on
any attribute

Content
management,
personalization,
mobile



In memory

Query by
key with microsecond
latency

Leaderboards,
real-time analytics,
caching



Graph

Quickly and easily
create and navigate
relationships
between data

Fraud detection,
social networking,
recommendation
engine



Time series

Collect, store, and
process data
sequenced by time

AWS IoT applications,
event tracking



Ledger

Complete,
immutable, and
verifiable history
of all changes to
application data

Systems of record,
supply chain,
healthcare,
registrations,
financial

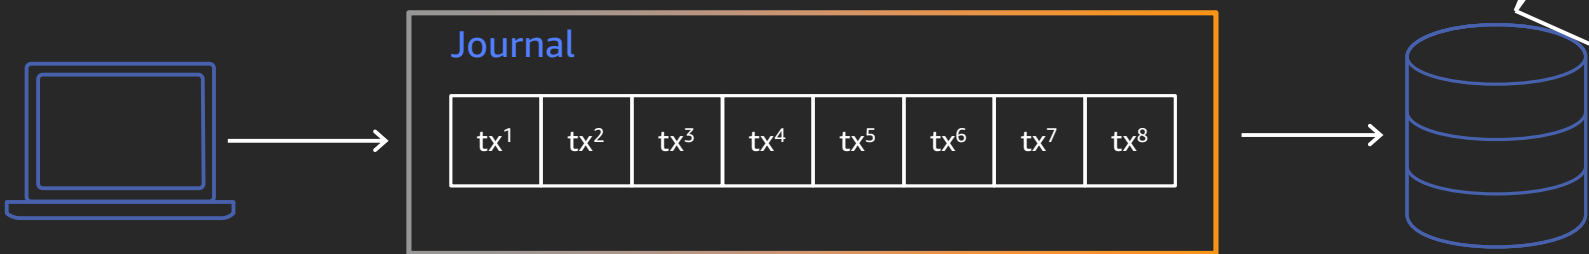
Traditional database architecture: The log

- Typically an internal implementation
- Used for replicating data
- Difficult—or impossible—to directly access



Amazon QLDB: The journal is the database

- QLDB’s journal has structural similarity to a database log
- All writes go to the journal—the journal determines state
- Journal handles concurrency, sequencing, cryptographic verifiability, and availability
- Accessible history of all transactions, document versions, document metadata



User #standard user data, the default

*select * from cars*

Manufacturer	Model	Year	VIN	Owner
Tesla	Model S	2012	123456789	Richard Roe

Committed #includes metadata

*select * from _ql_committed_cars*

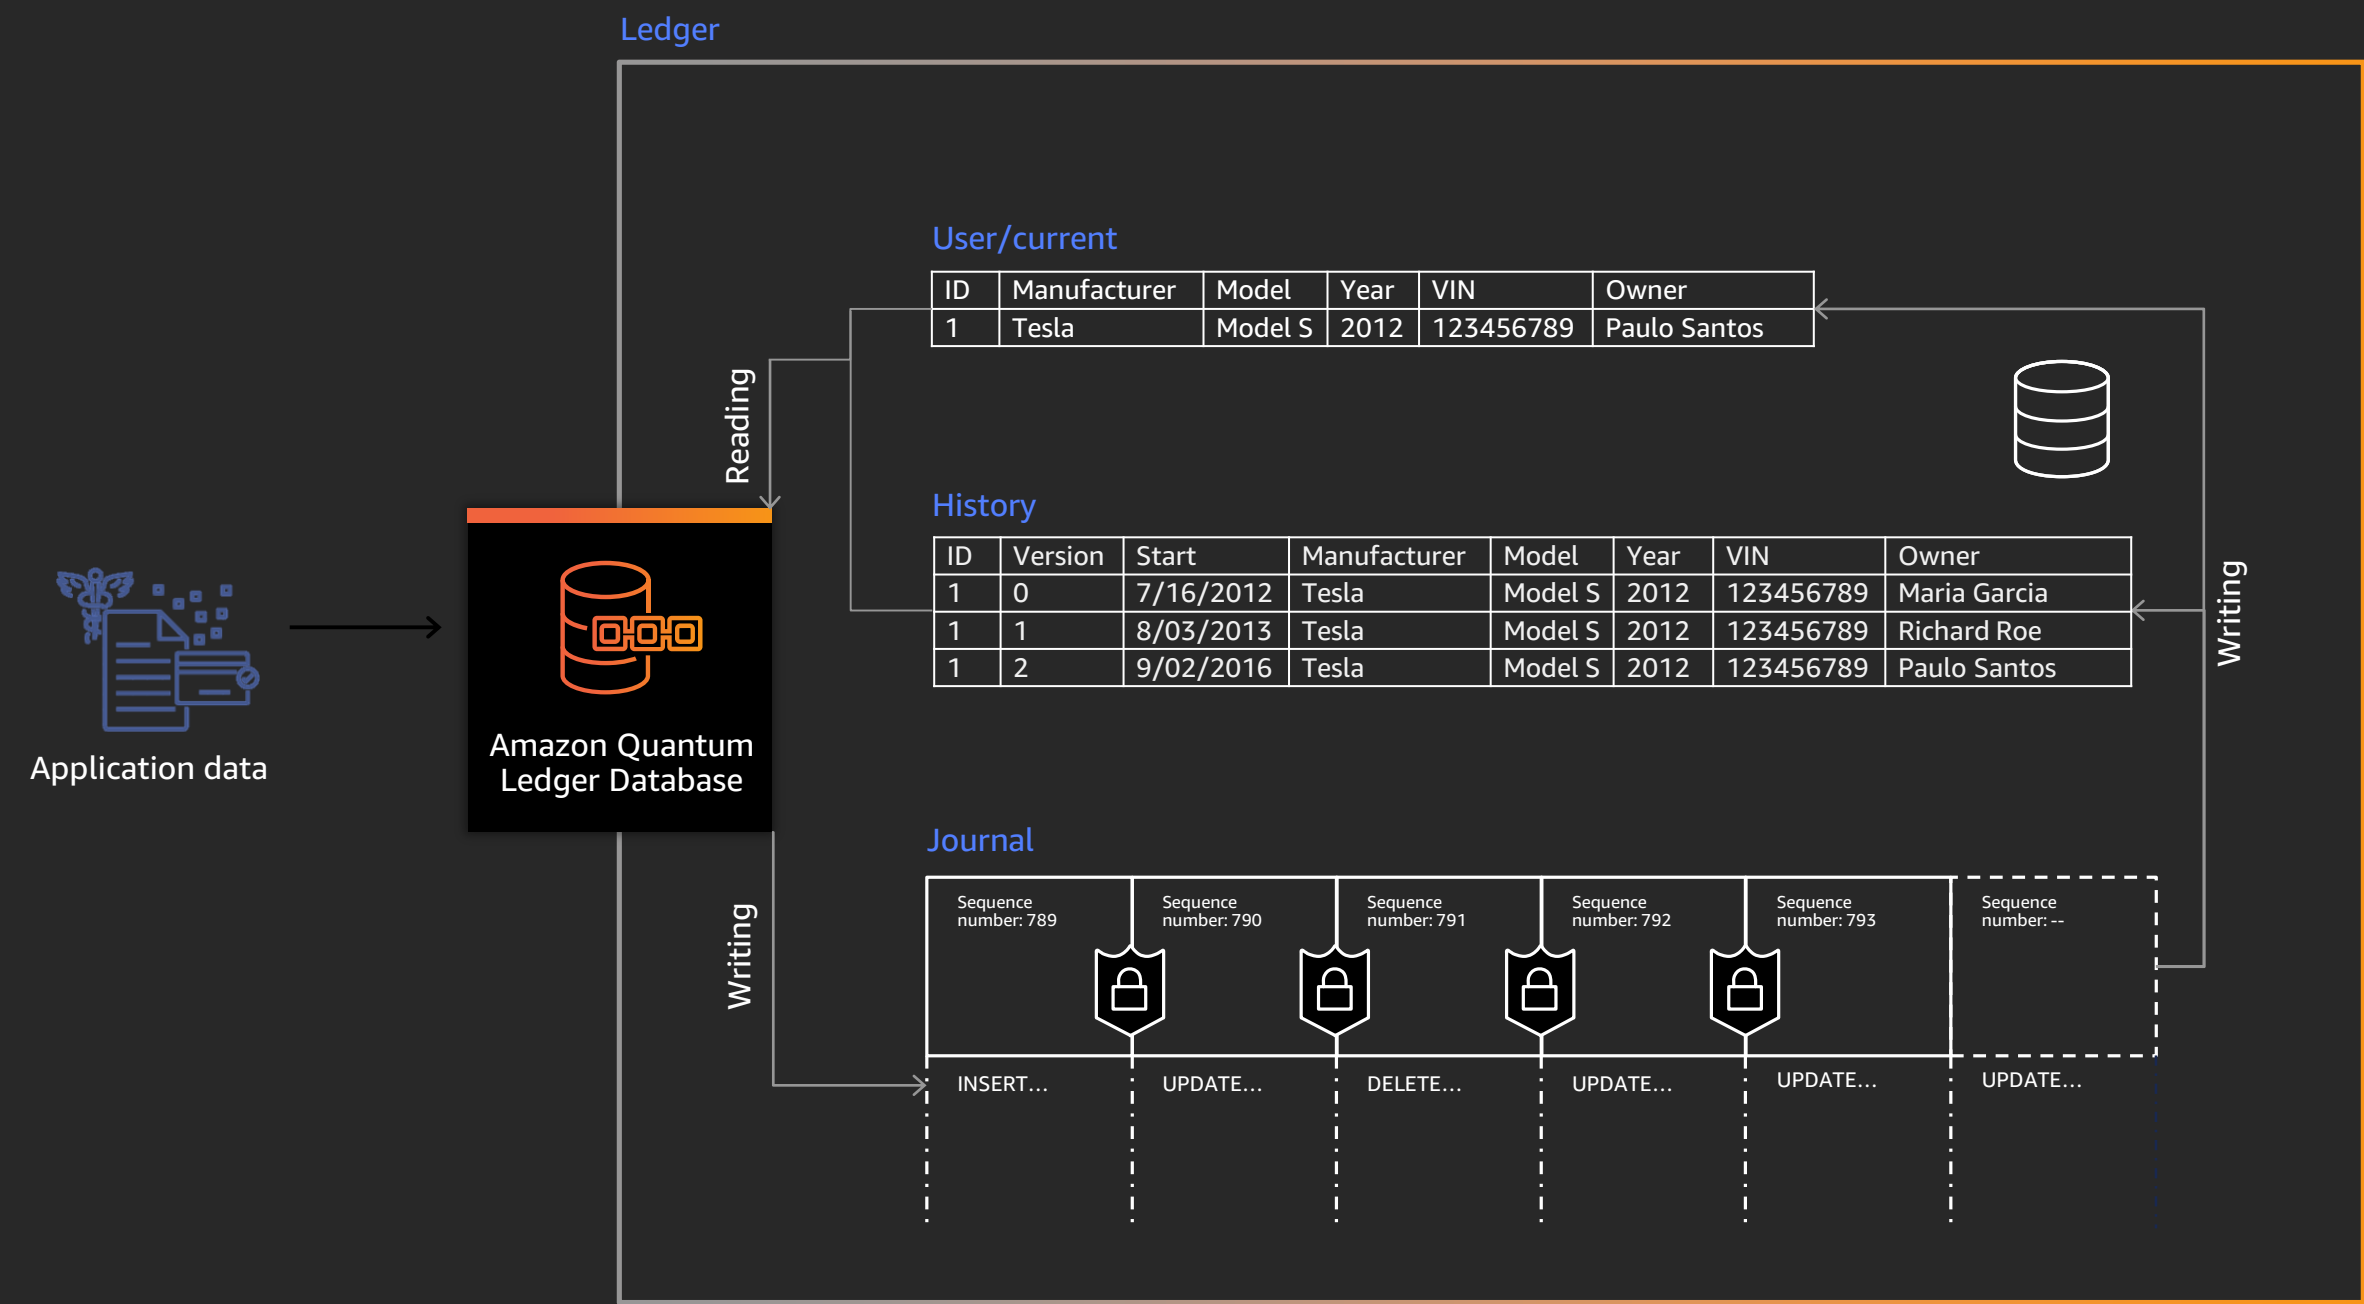
blockAddress	hash	data	metadata
{strandId:"liPT4AnNeWpE6YdoFiAK0U",sequenceNo:139}	{{1fWeNgIm59pgnkdBol83exdL+t/BAvVfICrX8N9cYN4=}}	{Manufacturer:"Tesla",Model:"Model S",Year:2012,VIN:123456789,Owner:"Richard Roe"}	{id:"Cxlp6cA36YOHylSChJ7dzW",version:1,txTime:2019-10-28T12:20:25.979Z,txId:"1nmeGu0dy4nl3KHcsfC4tr"}

history() #function to query document history

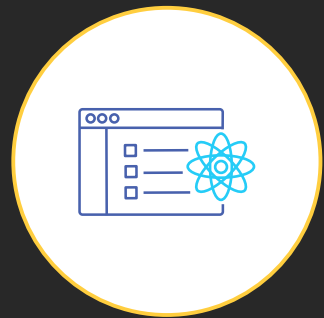
*SELECT * FROM history(cars) WHERE:*

blockAddress	hash	data	metadata
{strandId:"liPT4AnNeWpE6YdoFiAK0U",sequenceNo:136}	{{s6ytAZivsX2ukJNSwBVvYZ5KjnFQRJHRDrU0UKbYVRY=}}	{Manufacturer:"Tesla",Model:"Model S",Year:2012,VIN:123456789,Owner:"Maria Garcia"}	{id:"Cxlp6cA36YOHylSChJ7dzW",version:0,txTime:2019-10-28T12:19:50.098Z,txId:"JUJgLL39zof3a0LKTXNRgX"}
{strandId:"liPT4AnNeWpE6YdoFiAK0U",sequenceNo:139}	{{1fWeNgIm59pgnkdBol83exdL+t/BAvVfICrX8N9cYN4=}}	{Manufacturer:"Tesla",Model:"Model S",Year:2012,VIN:123456789,Owner:"Richard Roe"}	{id:"Cxlp6cA36YOHylSChJ7dzW",version:1,txTime:2019-10-28T12:20:25.979Z,txId:"1nmeGu0dy4nl3KHcsfC4tr"}

Amazon QLDB—Journal-first database



Transactions (ACID)



Highest to lowest

Isolation level	Potential issues
Serializable	—
Snapshot isolation	Potential write skew
Repeatable read	Phantom reads
Read committed	Phantom reads/non-repeatable reads
Read uncommitted	Phantom reads/non-repeatable reads/dirty reads

- Amazon QLDB supports the highest level of isolation
- There is no other mode for QLDB
- There is no risk that you'll see phantom reads, write skew, dirty reads, or other issues

Easy to use—Amazon Ion & PartiQL



Amazon Ion

```
/* Ion supports comments.
This is a "Vehicle" document */
{
  'VIN' : 'KM8SRDHF6EU074761' ,
  'MfgDate': `2017-03-01T` ,
  'Type':    'Truck' ,
  'Mfgr':    'Ford' ,
  'Model':   'F150'
  'Color':   'Black' ,
  'Specs': {
    'EngSize' : 3.3 ,(decimal)
    'Curbweight': 4878 , (int)
    'HP': 327 , (int)
    'BatterySize' : NULL.int ,
  }
}
```

PartiQL

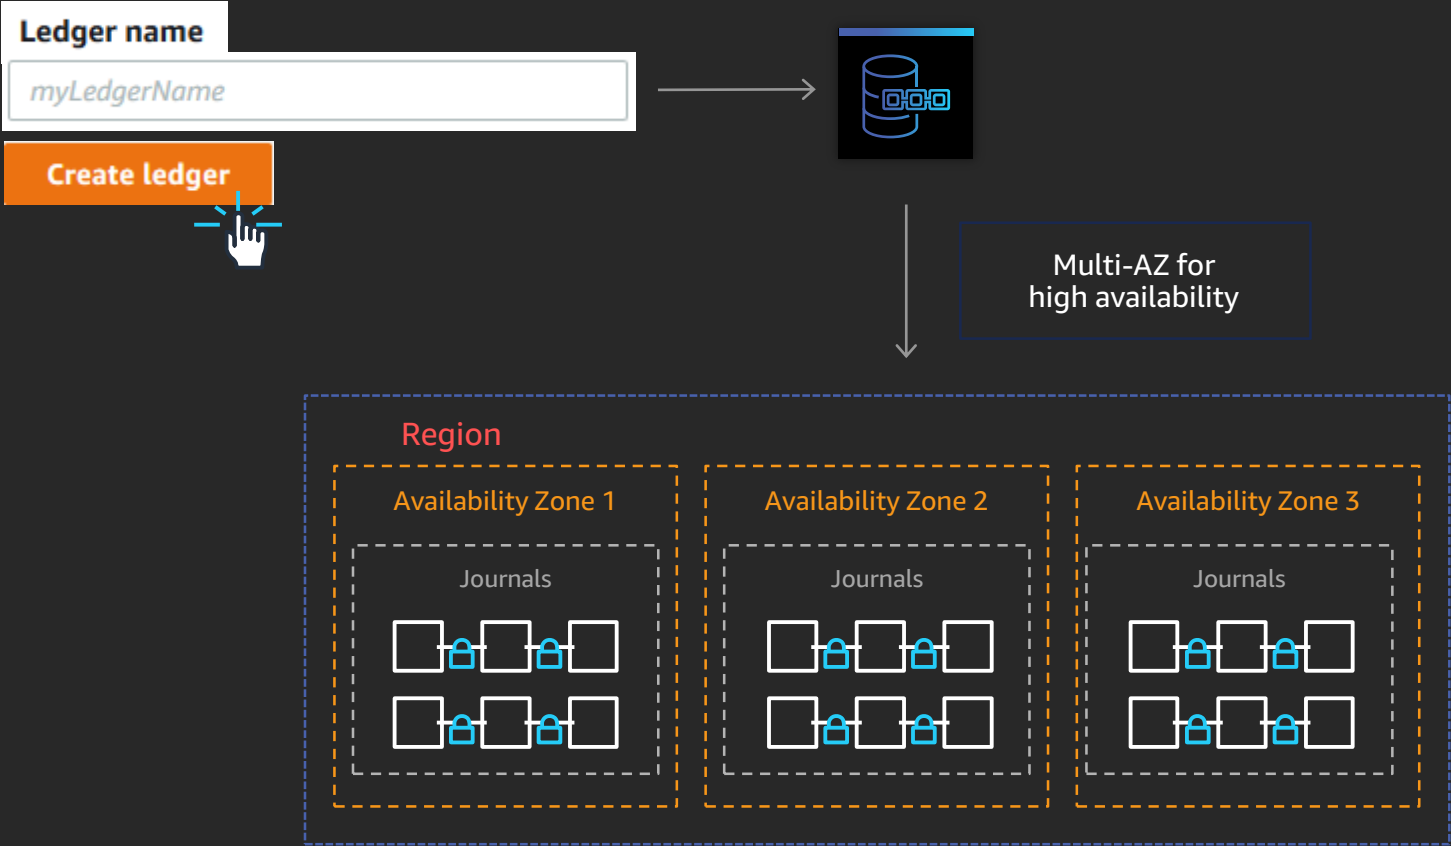
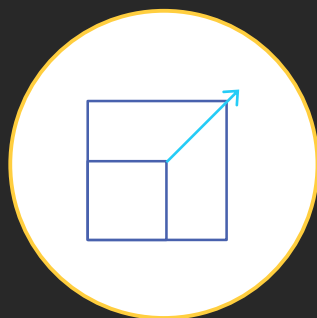
```
INSERT INTO cars
  { 'Manufacturer': 'Tesla',
    'Model': 'Model S',
    'Year': 2012,
    'VIN': 123456789,
    'Owner': 'Maria Garcia'
  }

UPDATE cars SET owner = 'Richard Roe'

WHERE VIN = 123456789

SELECT * FROM cars
```

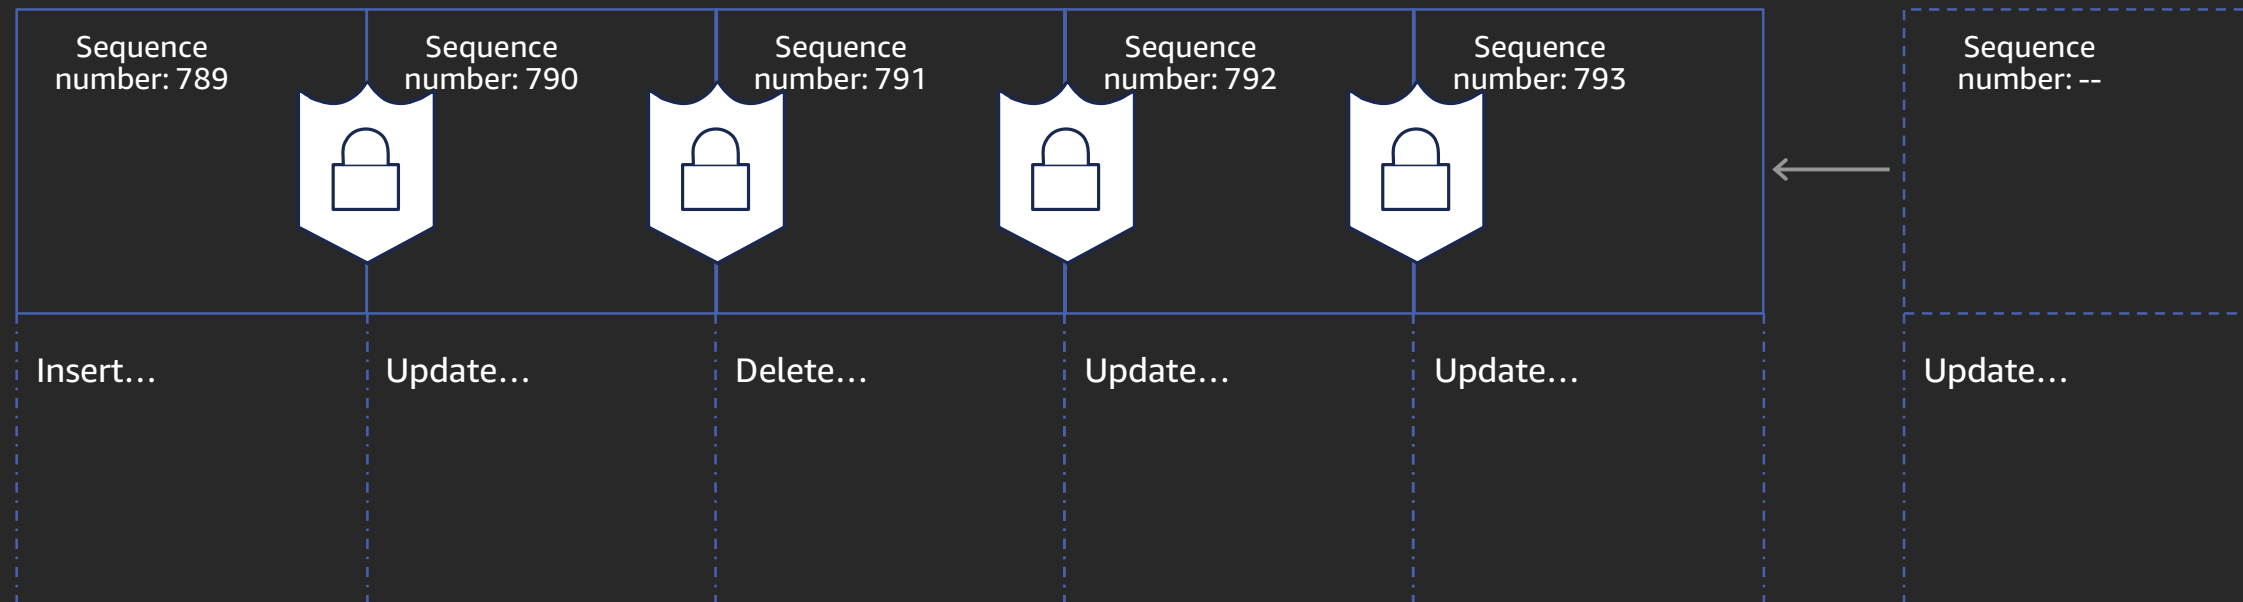
Serverless, scalable, highly available



Multiple copies per AZ
providing strong
durability

Immutable

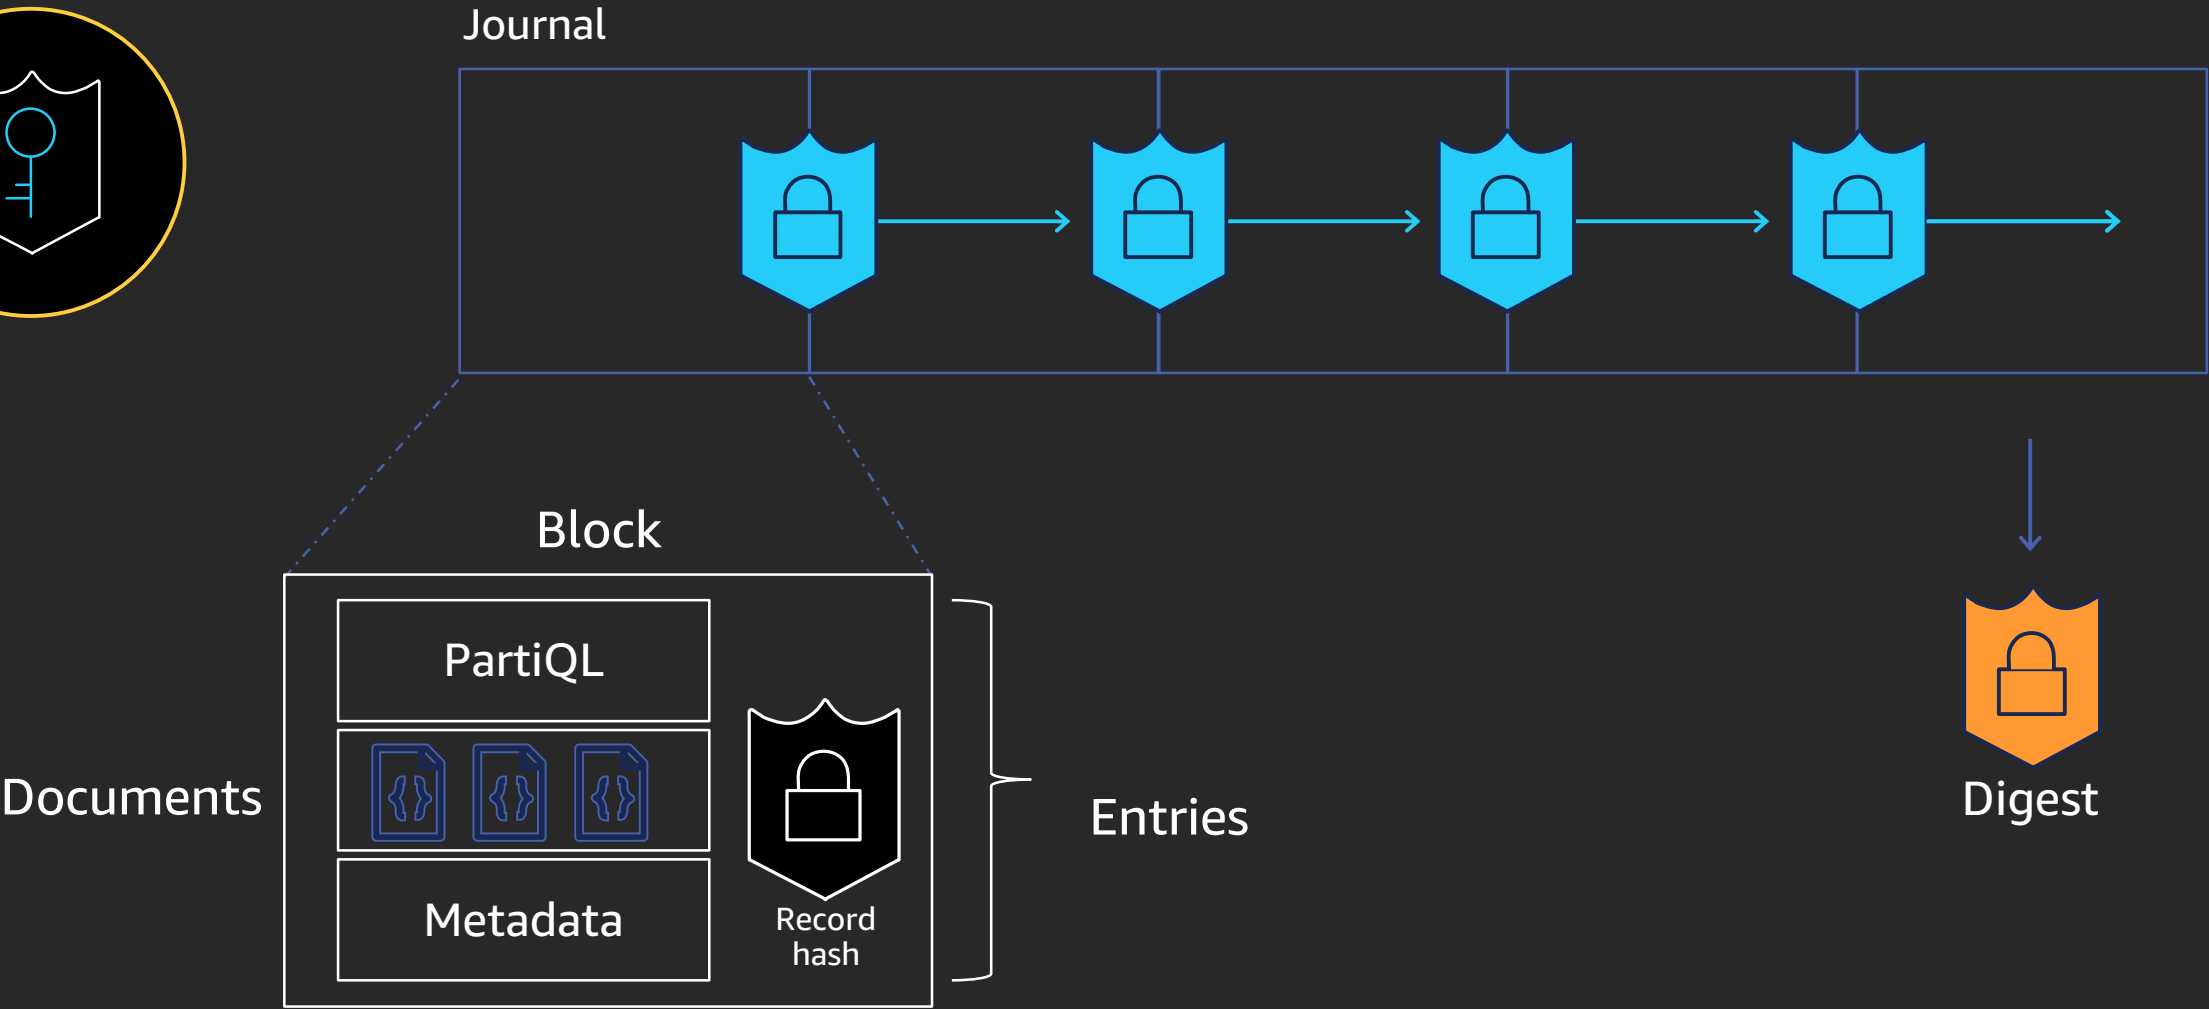
Records cannot be altered



- The journal is append only and sequenced
- There is no API or other method to alter committed data
- All operations, including deletes, are written to the journal

Cryptographic verification

Hash chaining using sha-256



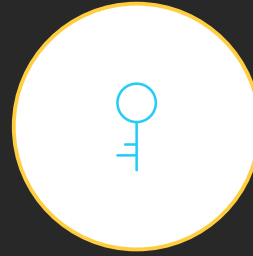
Amazon QLDB summary

Immutable



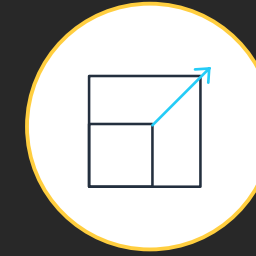
Append only, sequenced

Cryptographically verifiable



Hash chaining provides data integrity

Highly scalable



Serverless, highly available

Easy to use



Familiar SQL operators

ACID transactions



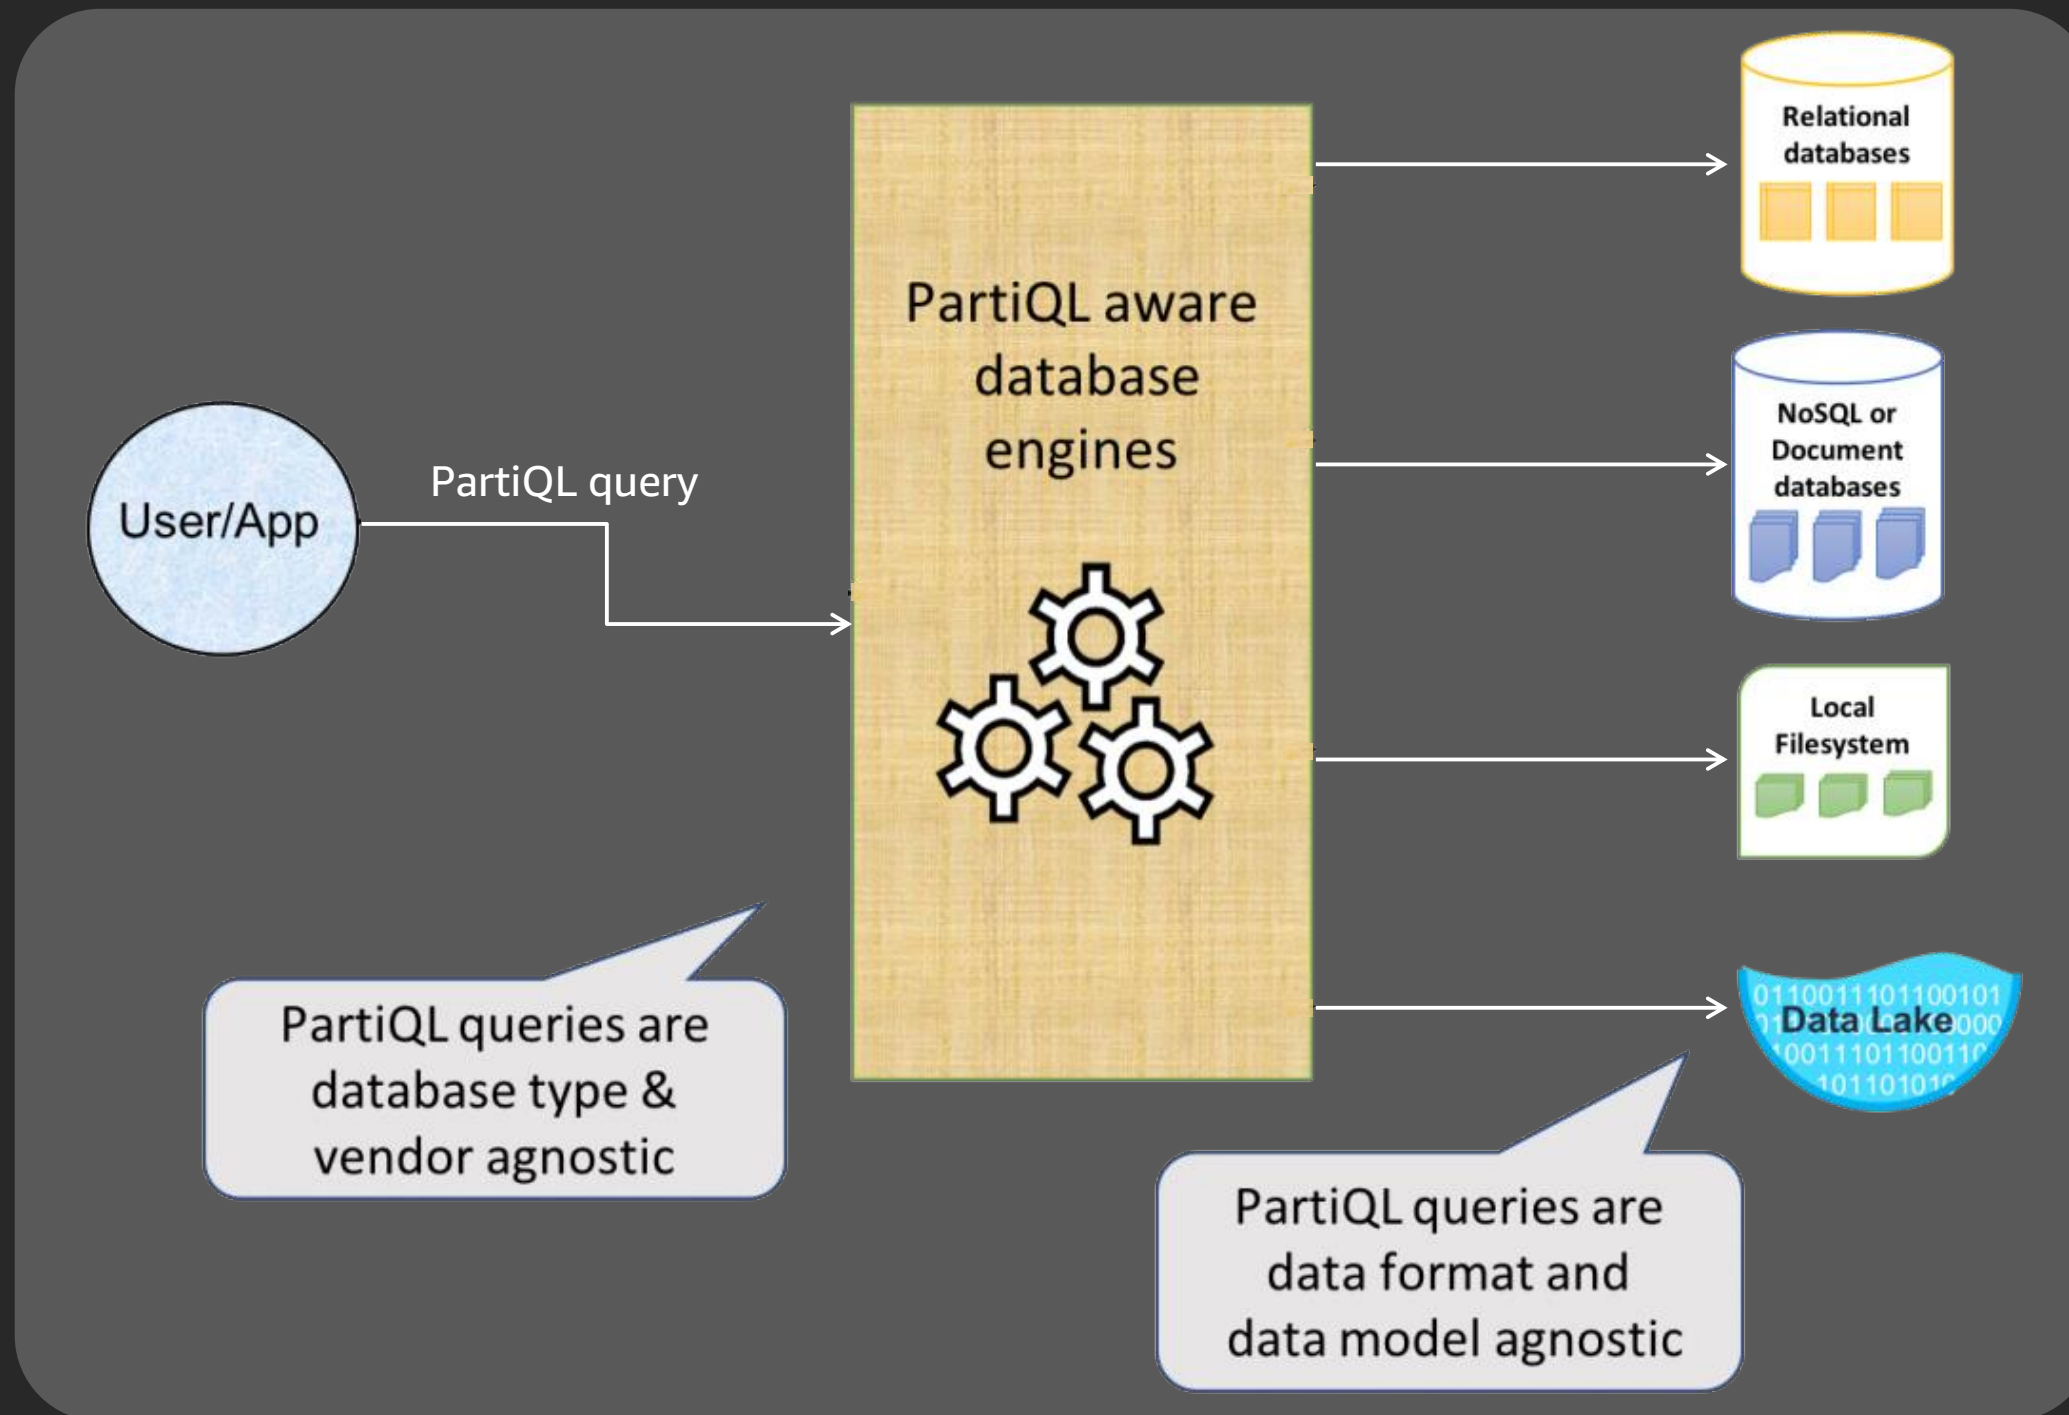
Fully serializable isolation

Journal first



The journal is the database

PartiQL—one query language for all your data



PartiQL—fundamental design tenets

1

SQL compatibility

2

First-class nested data

3

Minimal extensions

4

Format independence

5

Data store independence

Basic PartiQL query

Input data:

```
1 {
2   VIN: "1N4AL11D75C109151",
3   LicensePlateNumber: "LEWISR261LL",
4   State: "WA",
5   City: "Seattle",
6   PendingPenaltyTicketAmount: 90.25,
7   ValidFromDate: 2017-08-21T,
8   ValidToDate: 2020-05-11T,
9   Owners: {
10    PrimaryOwner: {
11      PersonId: "C6XZswBSmTS0myD5MW34B9"
12    },
13    SecondaryOwners: [
14      {PersonId: "5Ufgdlnj06gF5CWcOIu64s"},
15      {PersonId: "HrBDqUrWTHD3lN12CJ605q"}
16    ]
17  }
18 }
```

Query:

```
1 SELECT
2   s.VIN AS VIN
3 FROM
4   VehicleRegistration AS s
5 WHERE
6   s.City = 'Seattle'
```

Output:

```
1 {
2   VIN: "1N4AL11D75C109151"
3 }
```

Nested collections

Input data:

```
1 {
2   VIN: "1N4AL11D75C109151",
3   LicensePlateNumber: "LEWISR261LL",
4   State: "WA",
5   City: "Seattle",
6   PendingPenaltyTicketAmount: 90.25,
7   ValidFromDate: 2017-08-21T,
8   ValidToDate: 2020-05-11T,
9   Owners: {
10     PrimaryOwner: {
11       PersonId: "C6XZswBSmTS0myD5MW34B9"
12     },
13     SecondaryOwners: [
14       {PersonId: "5Ufgdlnj06gF5CWcOIu64s"},
15       {PersonId: "HrBDqUrWTHD3lN12CJ605q"}
16     ]
17   }
18 }
```

Query:

```
1 SELECT
2   s.VIN AS VIN,
3   p.PersonId as PersonId
4 FROM
5   VehicleRegistration AS s,
6   s.Owners AS o,
7   o.SecondaryOwners as p
8 WHERE
9   p.PersonId = '5Ufgdlnj06gF5CWcOIu64s'
```

Output:

```
1 {
2   VIN: "1N4AL11D75C109151",
3   PersonId: "5Ufgdlnj06gF5CWcOIu64s"
4 }
```

Ion features

1

Value based

2

Supports primitive types and container types

3

Flexible and scalable

4

Hierarchy

5

Human readable

6

Efficient for machine

Ion types

Primitive types

bool, int, float, decimal,
timestamp, string,
symbol, blob, clob, null

Collection types

struct, list

- SQL Scalar types are covered by their Ion counterparts
- Ion's struct type is equivalent to an SQL tuple
- In addition to NULL, PartiQL has a MISSING type, meaning missing field

JSON versus Amazon Ion

JSON document

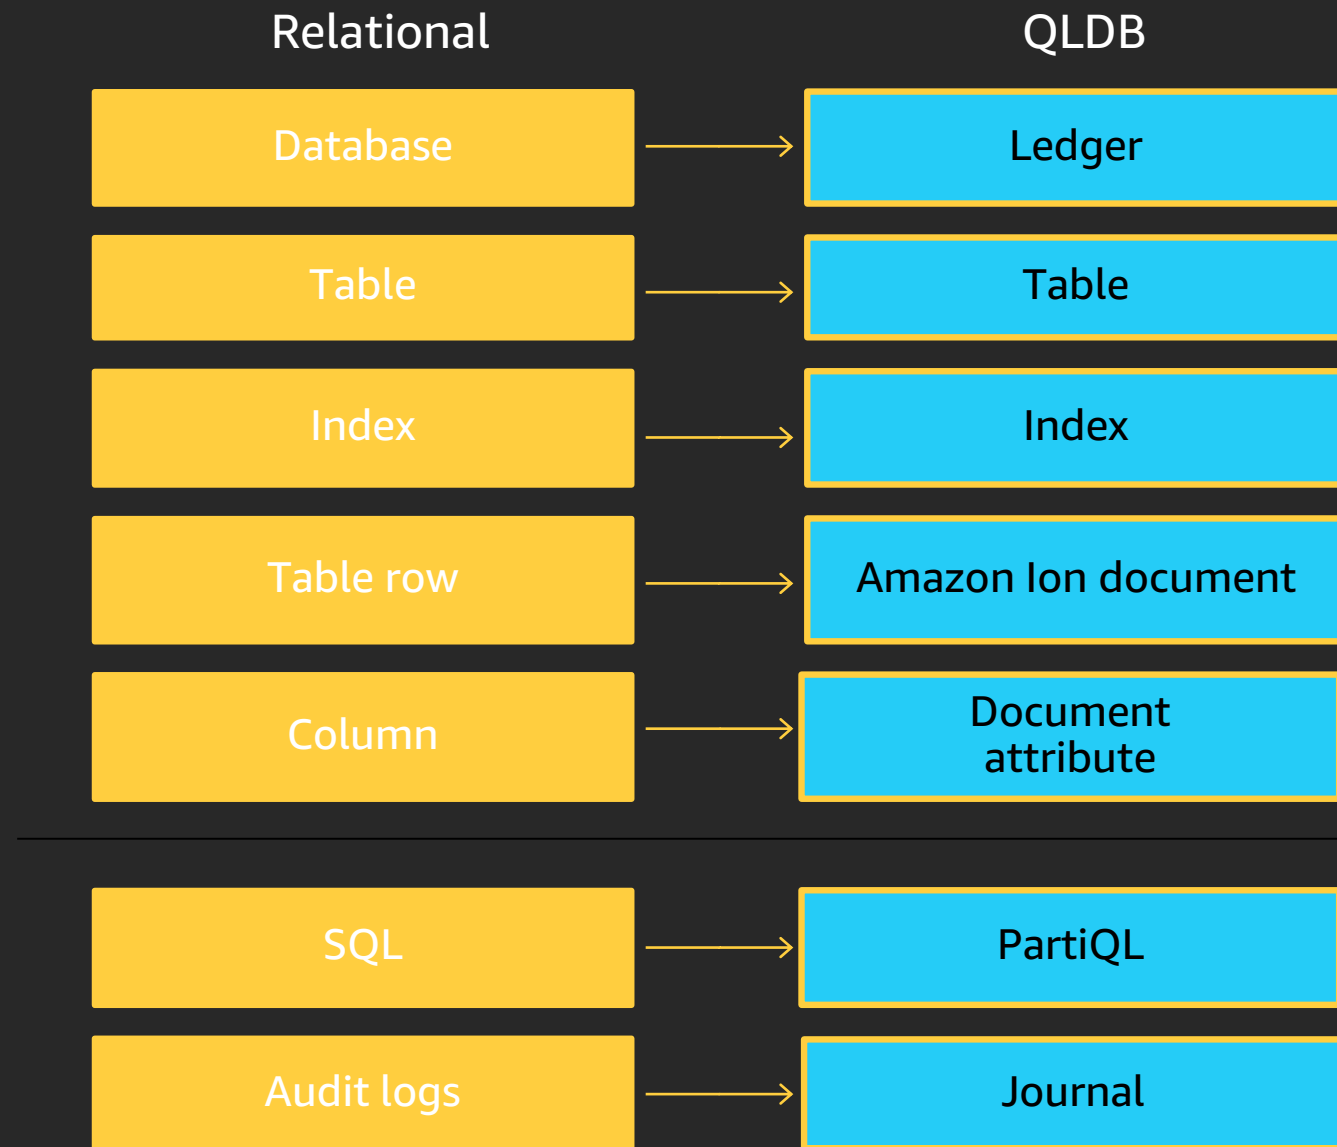
```
{
  "VIN" : "KM8SRDHF6EU074761",
  "MfgDate" : "2017-03-01",
  "Type" : "Truck",
  "Mfgr" : "Ford",
  "Model" : "F150",
  "Color" : "Black",
  "Specs" : {
    "EngSize" : 3.3,
    "CurbWeight": 4878,
    "HP": 327,
    "BatterySize": null
  }
}
```

Amazon Ion document

```
/* Ion supports comments. */
{
  VIN : "KM8SRDHF6EU074761", // (String)
  MfgDate : 2017-03-01T, // (timestamp)
  Type : "Truck",
  Mfgr : "Ford",
  Model : "F150",
  Color : "Black",
  Specs: {
    EngSize : 3.3 // (decimal)
    CurbWeight : 4878 // (int)
    HP : 327 // (int)
    BatterySize : null.int
  }
}
```

<https://github.com/amzn/ion-java>

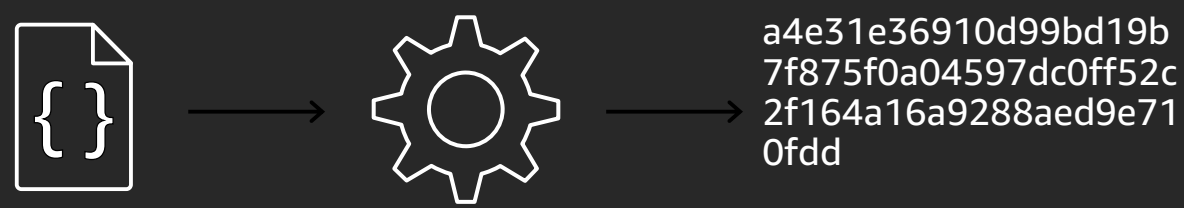
Mapping constructs between RDBMS & QLDB



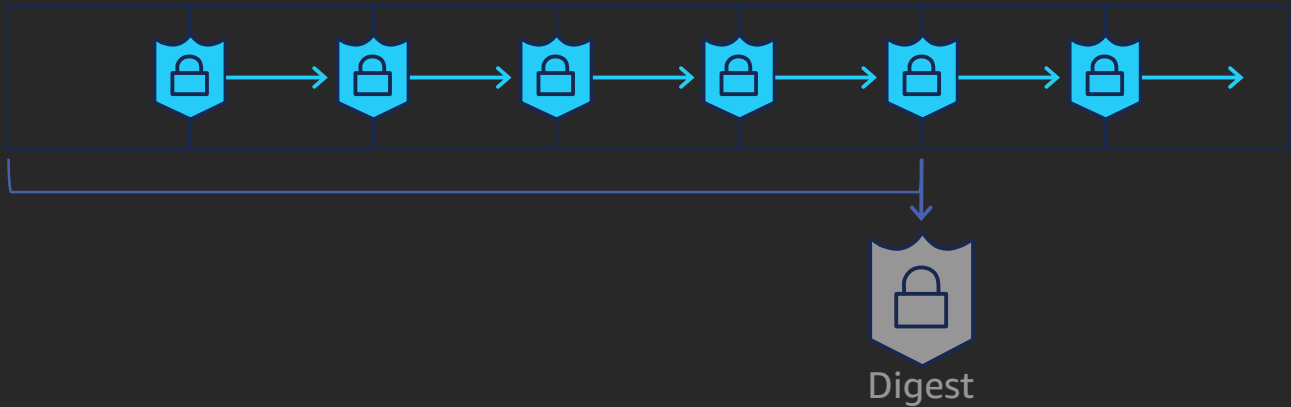
Cryptographic verifiability

Four basic steps to seeing how QLDB's verifiability works

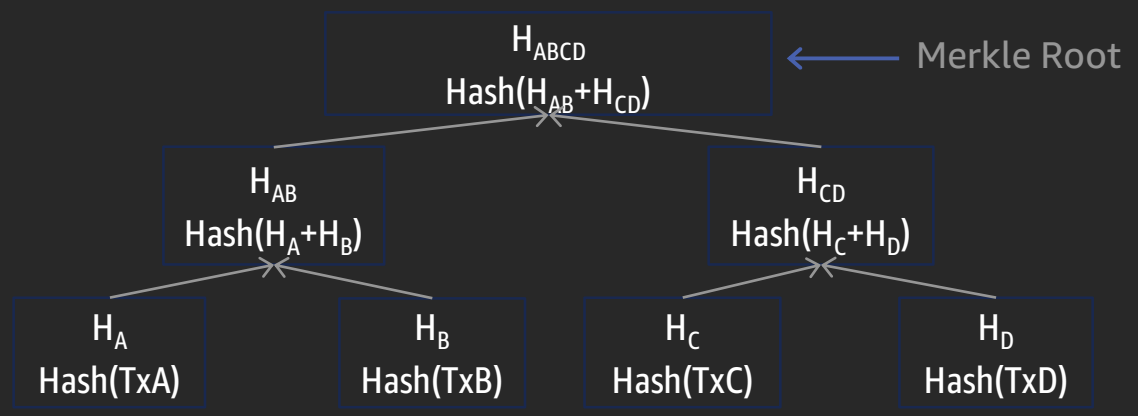
SHA256: Unique signature of a document



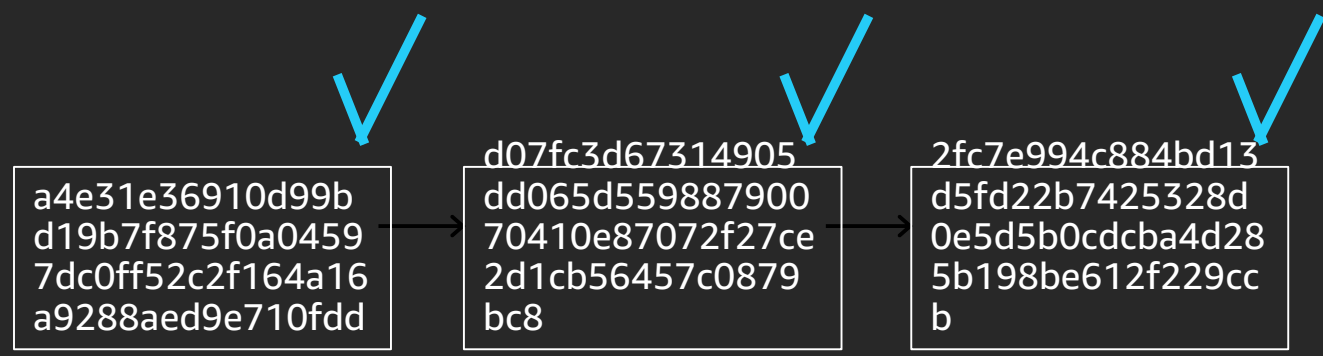
Digest: Periodic hash covering all history



Merkle trees: Chaining past hashes together



Proof: A chain of hashes linking a document to its digest



How it works

```
INSERT INTO cars
{ 'Manufacturer':'Tesla',
  'Model':'Model S',
  'Year':2012,
  'VIN':123456789,
  'Owner':'Maria Garcia' }
```

Cars

C

ID	Manufacturer	Model	Year	VIN	Owner

History

H

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner

Journal

J

```
INSERT cars
ID:1
Manufacturer: Tesla
Model: Model S
Year: 2012
VIN: 123456789
Owner: Maria Garcia

Metadata: {
Date:07/16/2012
}
```

H(T₁)

How it works

```
INSERT INTO cars
{ 'Manufacturer':'Tesla',
  'Model':'Model S',
  'Year':2012,
  'VIN':123456789,
  'Owner':'Maria Garcia' }
```

Cars

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Maria Garcia

History

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner
1	0	7/16/2012	Tesla	Model S	2012	123456789	Maria Garcia

Journal



How it works

```
UPDATE cars SET owner = 'Richard Roe' WHERE
VIN = 123456789
```

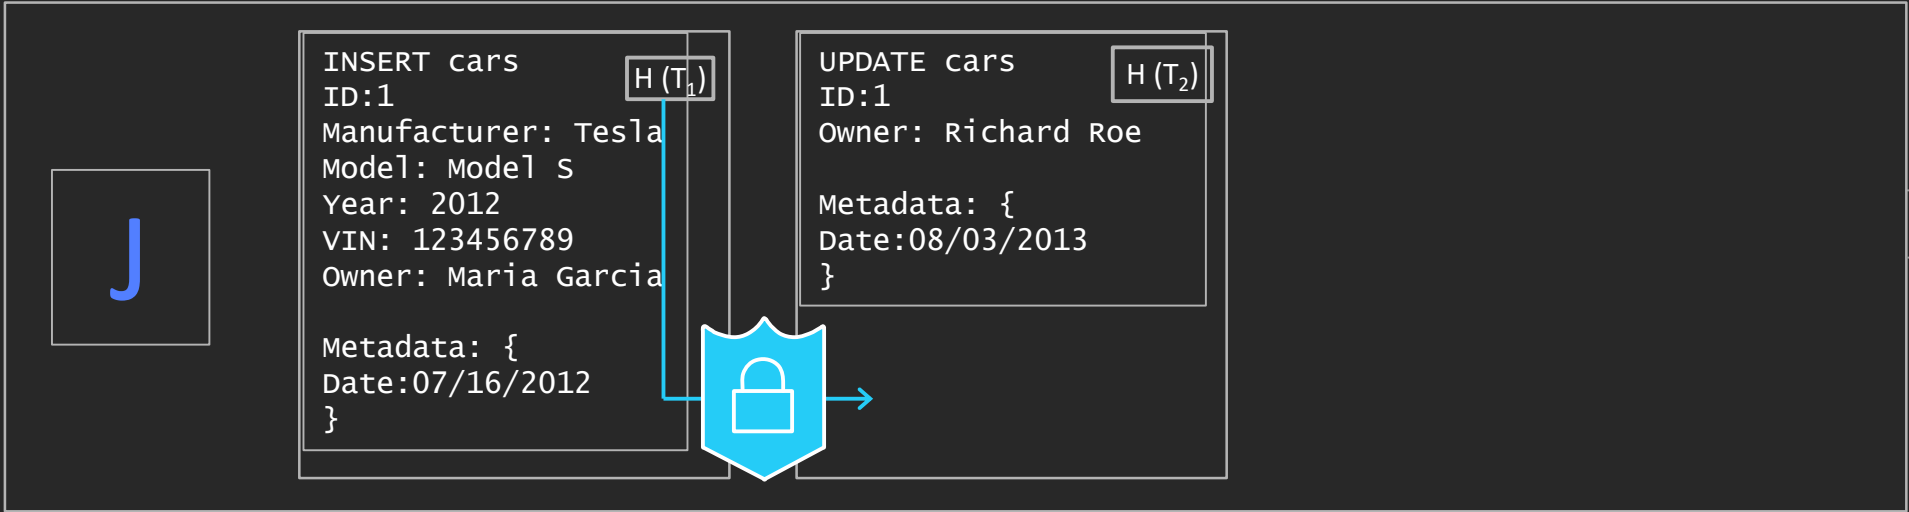
Cars

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Richard Roe

History

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner
1	0	7/16/2012	Tesla	Model S	2012	123456789	Maria Garcia
1	1	8/03/2013	Tesla	Model S	2012	123456789	Richard Roe

Journal



How it works

```
DELETE FROM cars WHERE VIN = 123456789
```

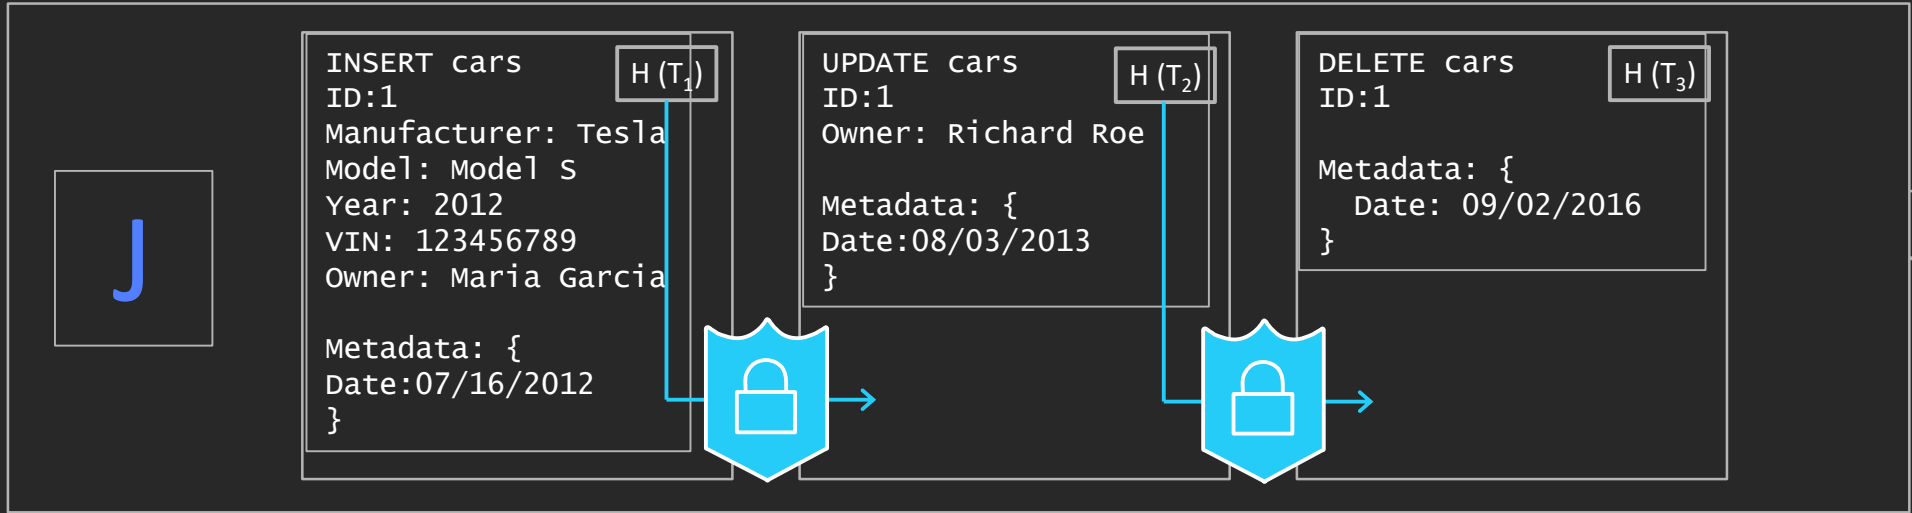
Cars

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Richard Roe

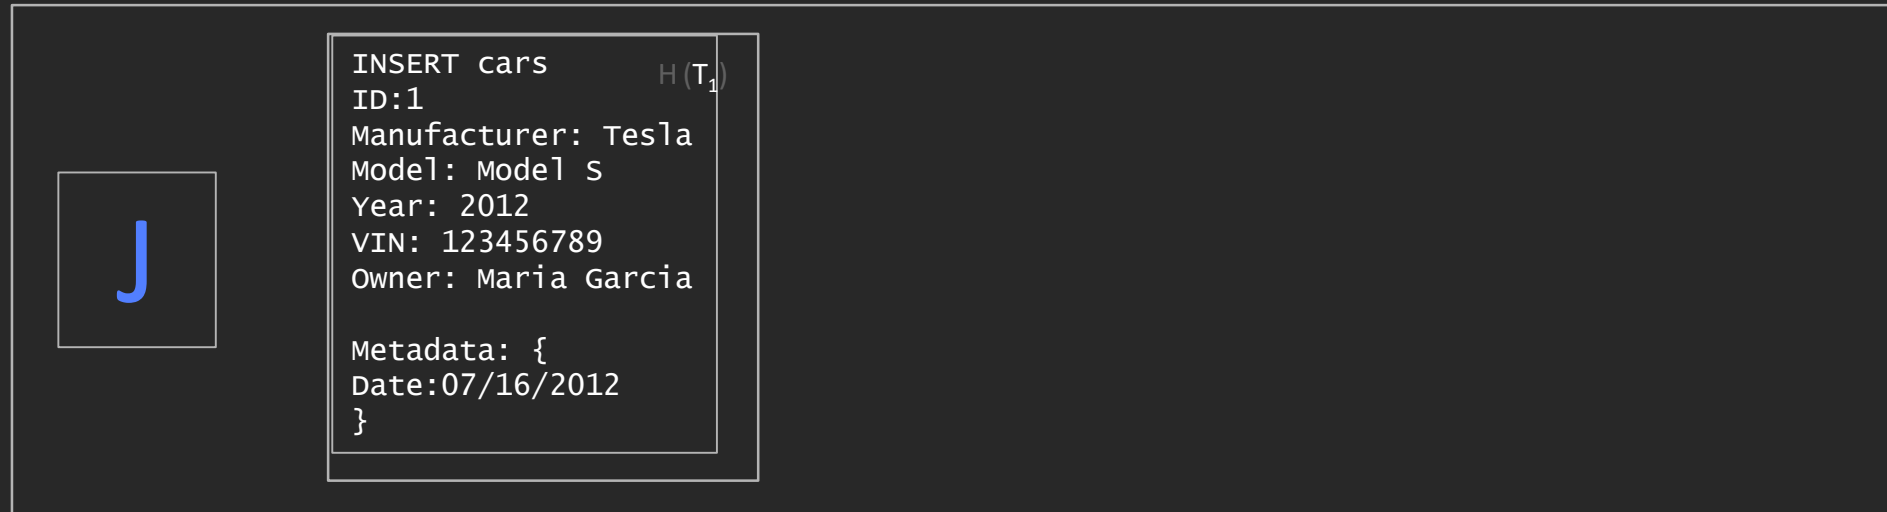
History

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner
1	0	7/16/2012	Tesla	Model S	2012	123456789	Maria Garcia
1	1	8/03/2013	Tesla	Model S	2012	123456789	Richard Roe
1	2	9/02/2016	Deleted				

Journal



Walk through a hash chain



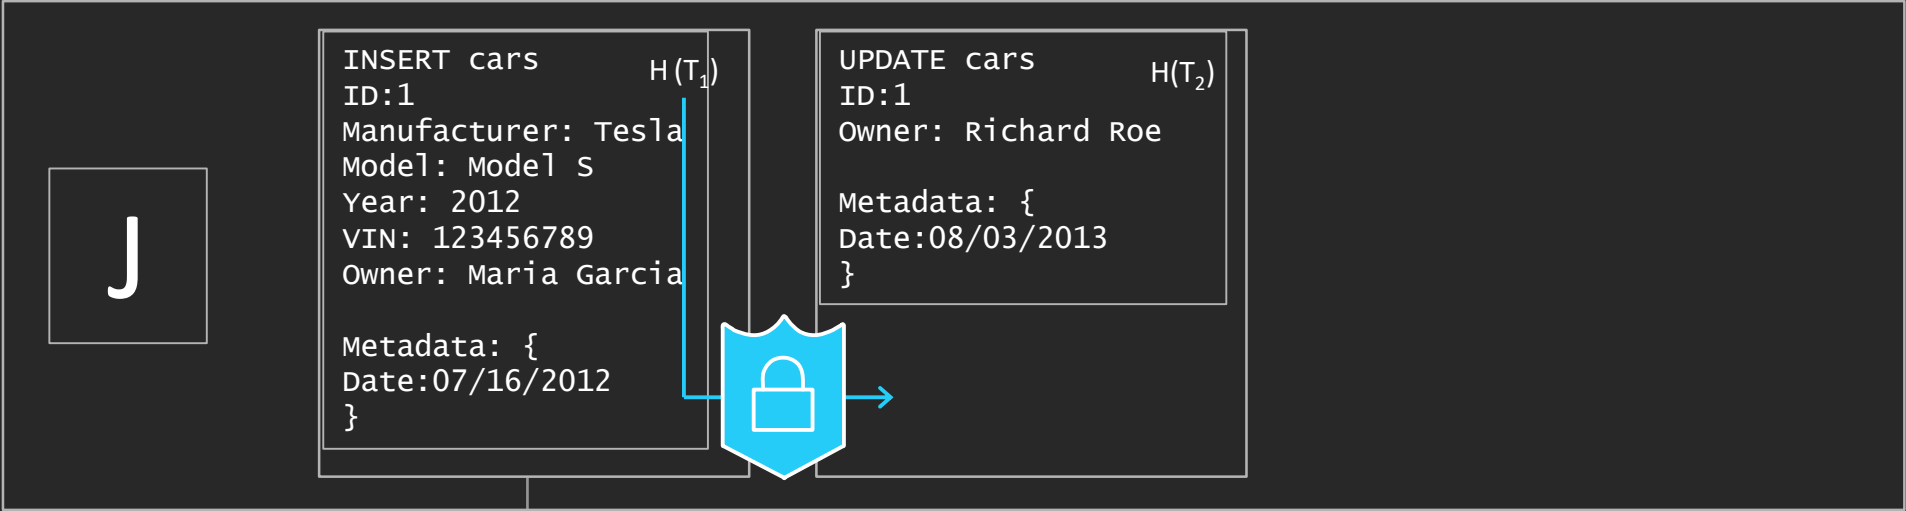
```
INSERT cars
ID:1
Manufacturer: Tesla
Model: Model S
Year: 2012
VIN: 123456789
Owner: Maria Garcia
```

```
Metadata: {
Date:07/16/2012
}
```

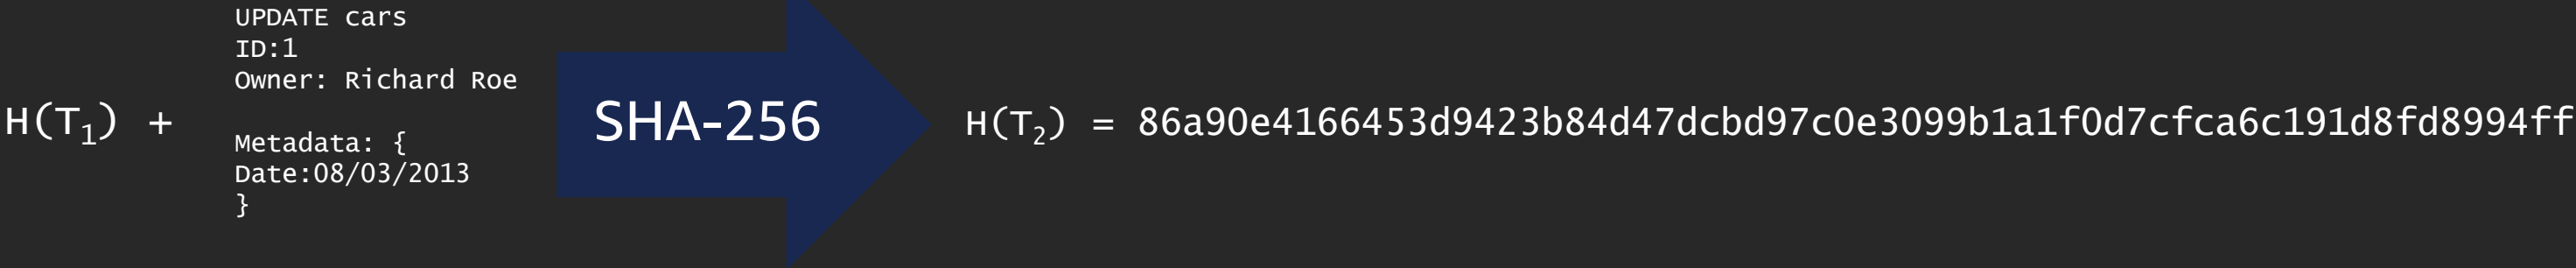
SHA-256

$H(T_1) = 2526f16306c819d651af075934170d2430d246d9ab98d975d28a83baded47ca7$

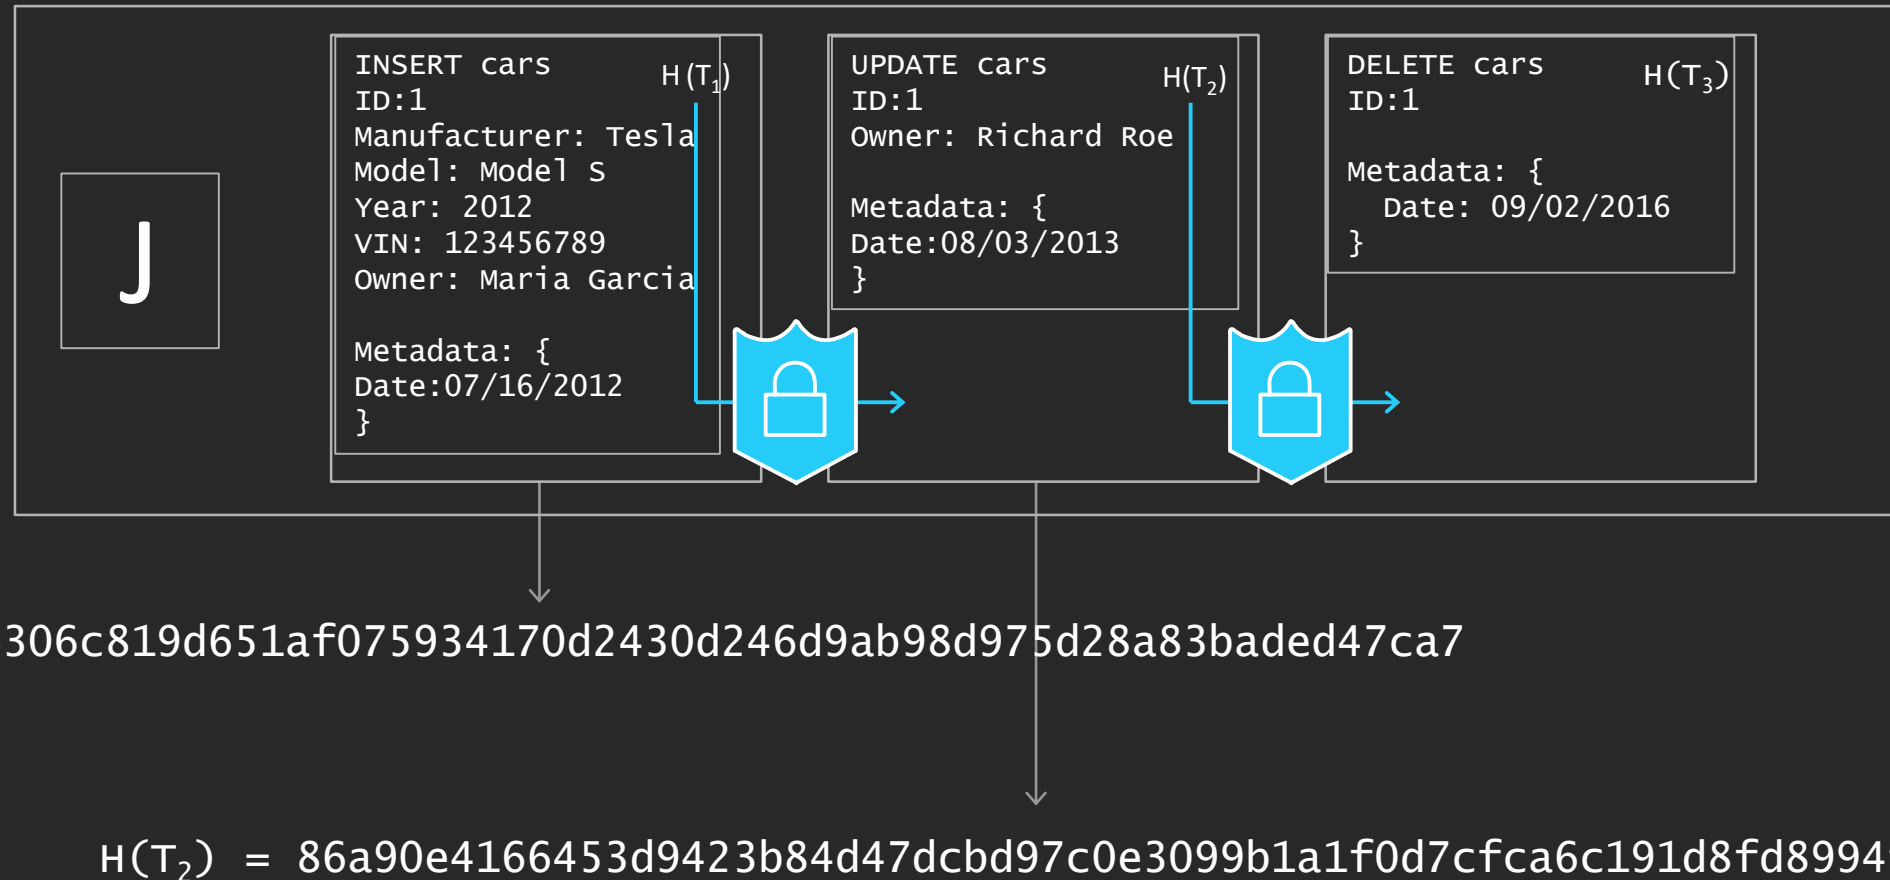
Hashing and chaining transactions



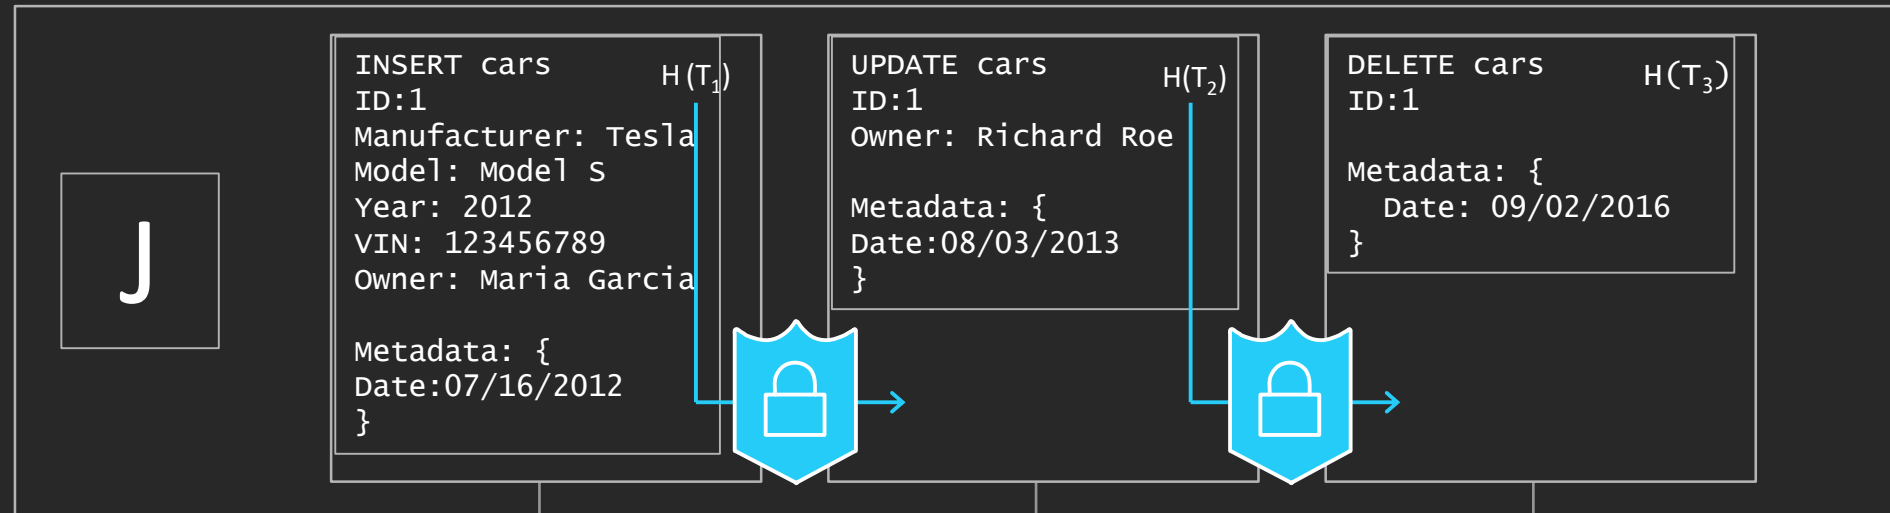
$H(T_1) = 2526f16306c819d651af075934170d2430d246d9ab98d975d28a83baded47ca7$



Hashing and chaining transactions



Hashing and chaining transactions

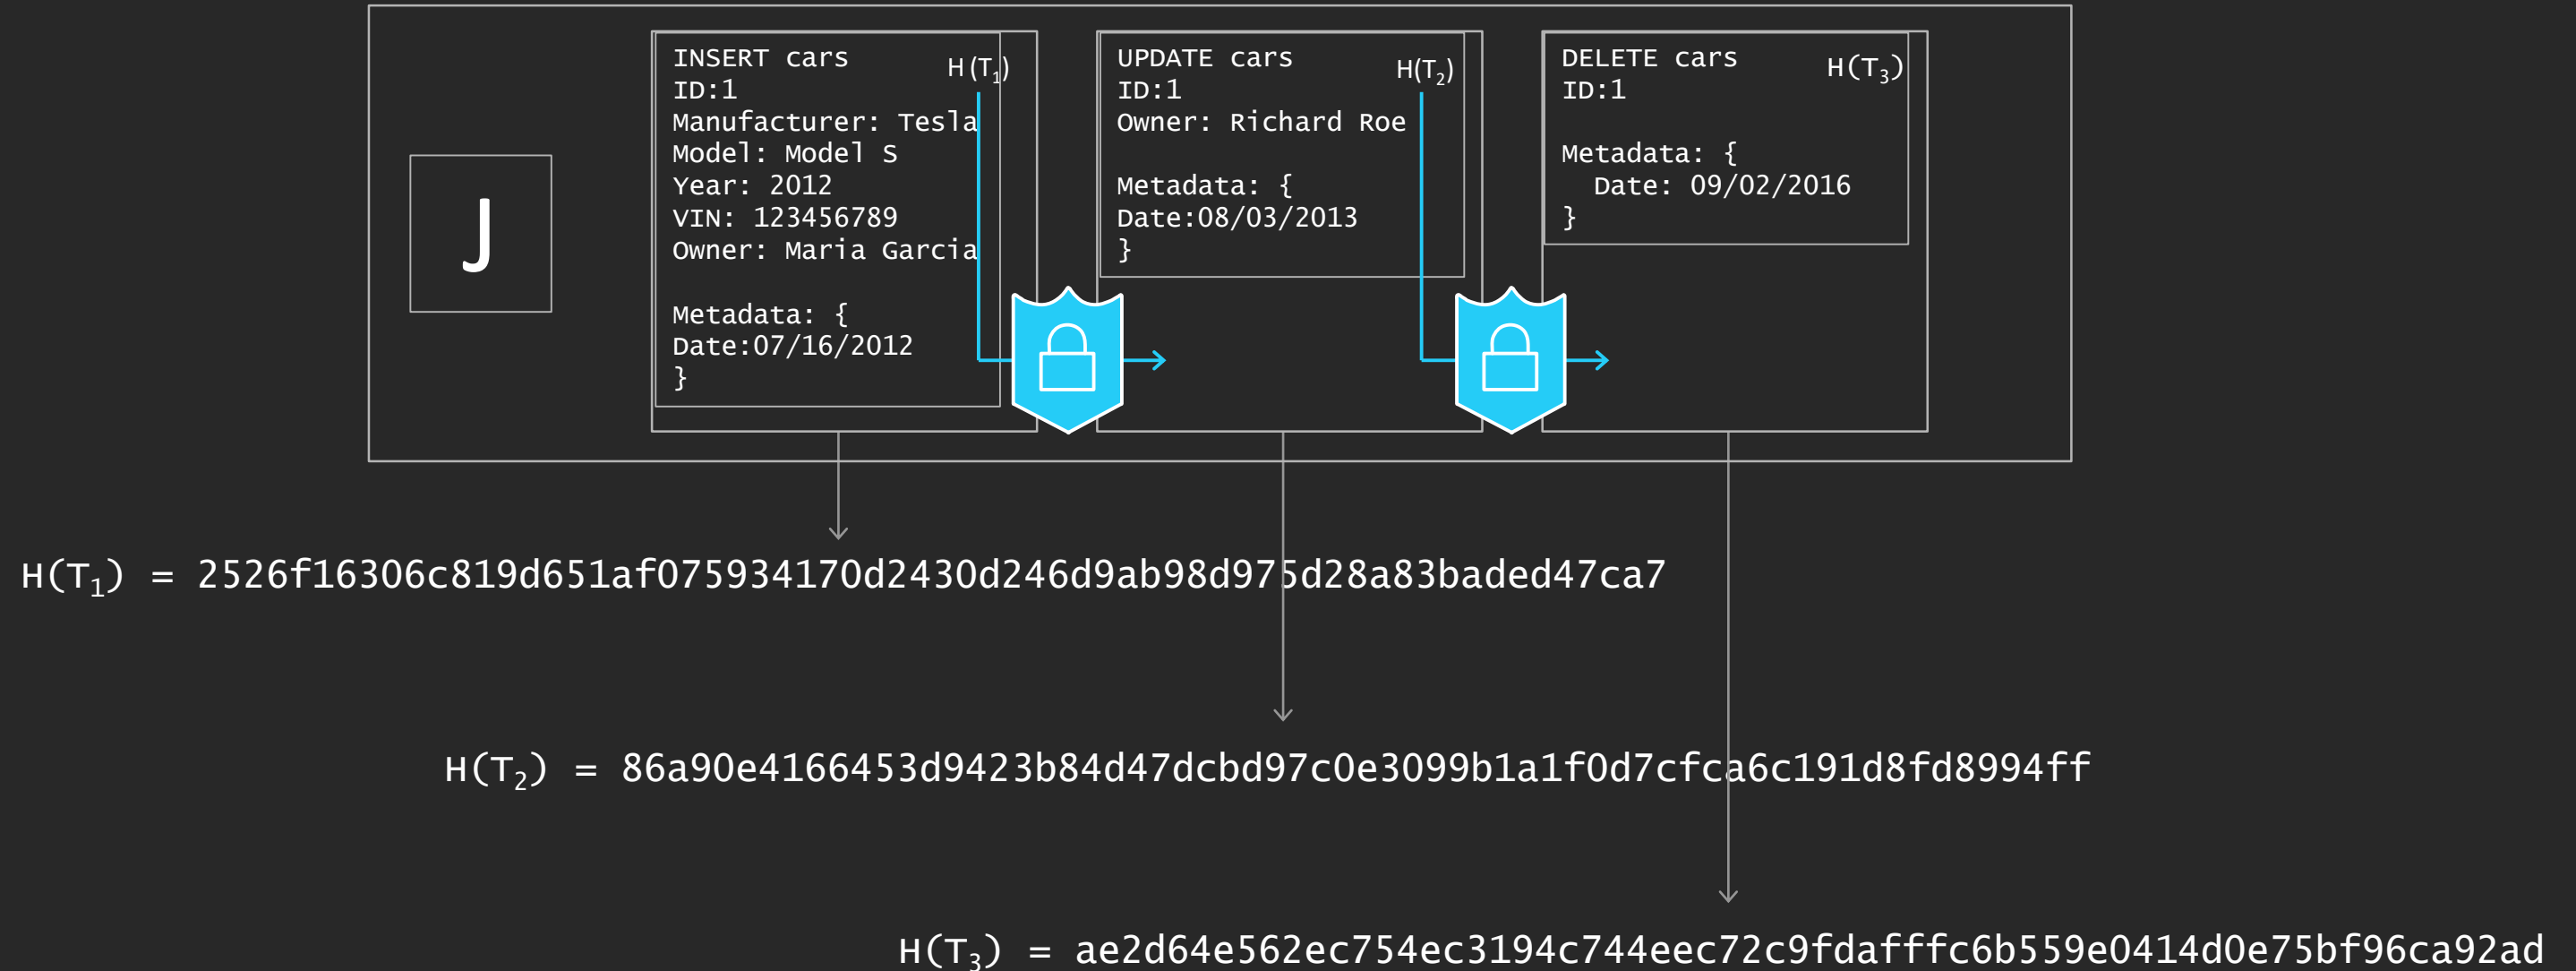


$H(T_1) = 2526f16306c819d651af075934170d2430d246d9ab98d975d28a83bade47ca7$

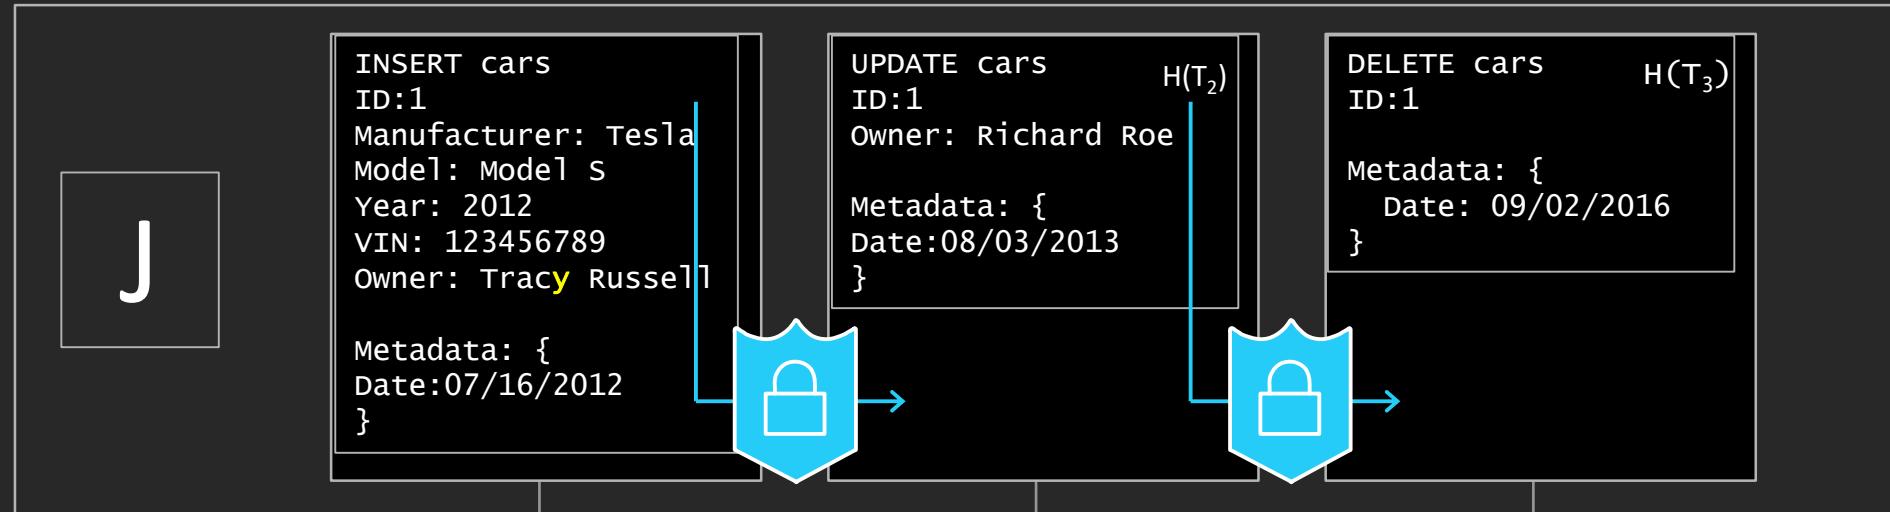
$H(T_2) = 86a90e4166453d9423b84d47dcbd97c0e3099b1a1f0d7cfca6c191d8fd8994ff$

$H(T_3) = ae2d64e562ec754ec3194c744eec72c9fdafffc6b559e0414d0e75bf96ca92ad$

A digest is a hash value at a point in time



Changing committed data breaks the chain



$H(T_1) =$ ~~2526f16306c819d651af075934170d2430d246d9ab98d975d28a83bade47ca7~~

$H(T_1) =$ 25d0b44e6e8878151646ffc1fea4eb85c3e4bf4baec212a9fcf67b6d5a81e01a

$H(T_2) =$ ~~86a90e4166453d9423b84d47dcbd97c0e3099b1a1f0d7cfca6c191d8fd8994ff~~

$H(T_2) =$ a90a9898c7e4b1aab19c705b554afd9e0bf6539bb0346df19be362ff63001098

$H(T_3) =$ ~~ae2d64e562ec754ec3194c744ecc72c9fdafffc6b559e0414d0e75bf96ca92ad~~

$H(T_3) =$ c6268578a24dbe0c7cfba07bd967411a35462b8c875d42f1991faad02c0ac93c

Challenges customers face

Building ledgers with traditional databases



Resource
intensive



Difficult to
manage and scale

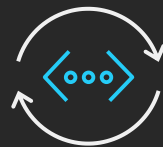


Error prone and
incomplete



Impossible
to verify

Blockchain approaches



Designed for a
different purpose



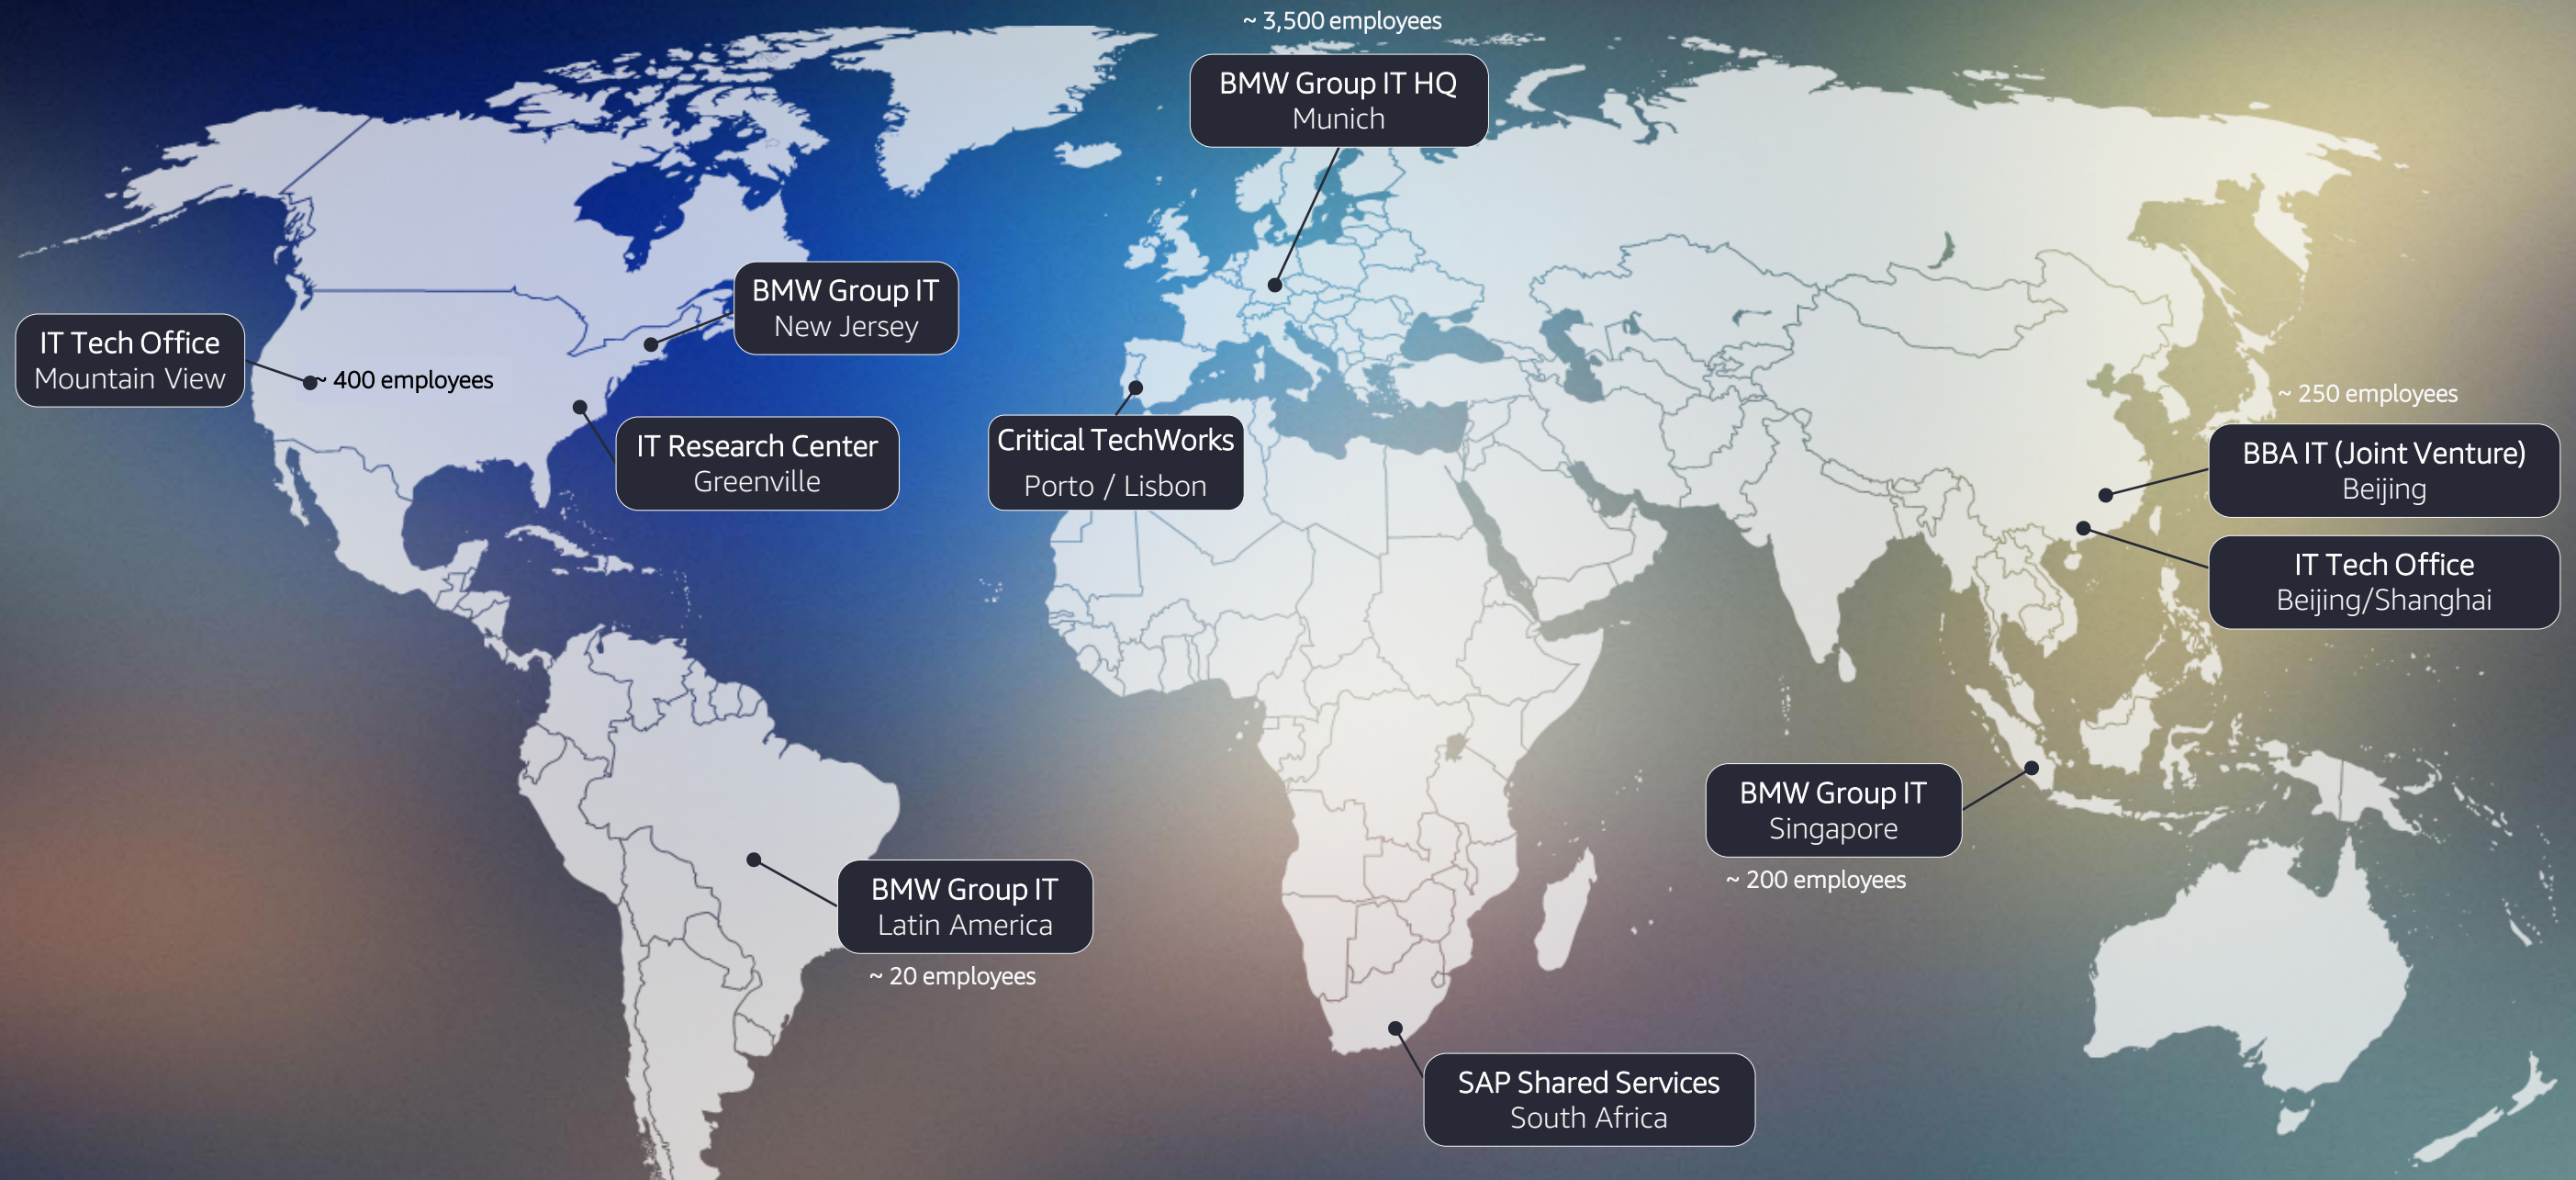
Adds unnecessary
complexity

Leveraging Amazon QLDB as a trusted verifiable ledger for automotive data


Emil Djerekarov

Technical Architect
BMW Group

BMW Group IT



Ledger technologies within automotive industry



Supply chain
& production



Vehicle

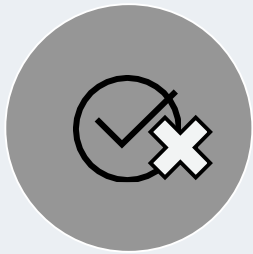


Artificial intelligence

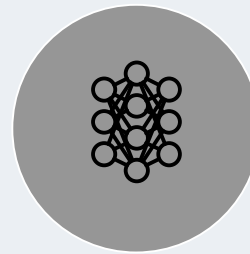


Verifiable data provenance

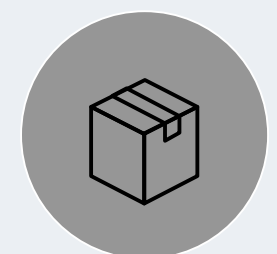
Ledger technologies enable new business models



Data must be
verified and
auditable



Multiple organizations
and partners involved

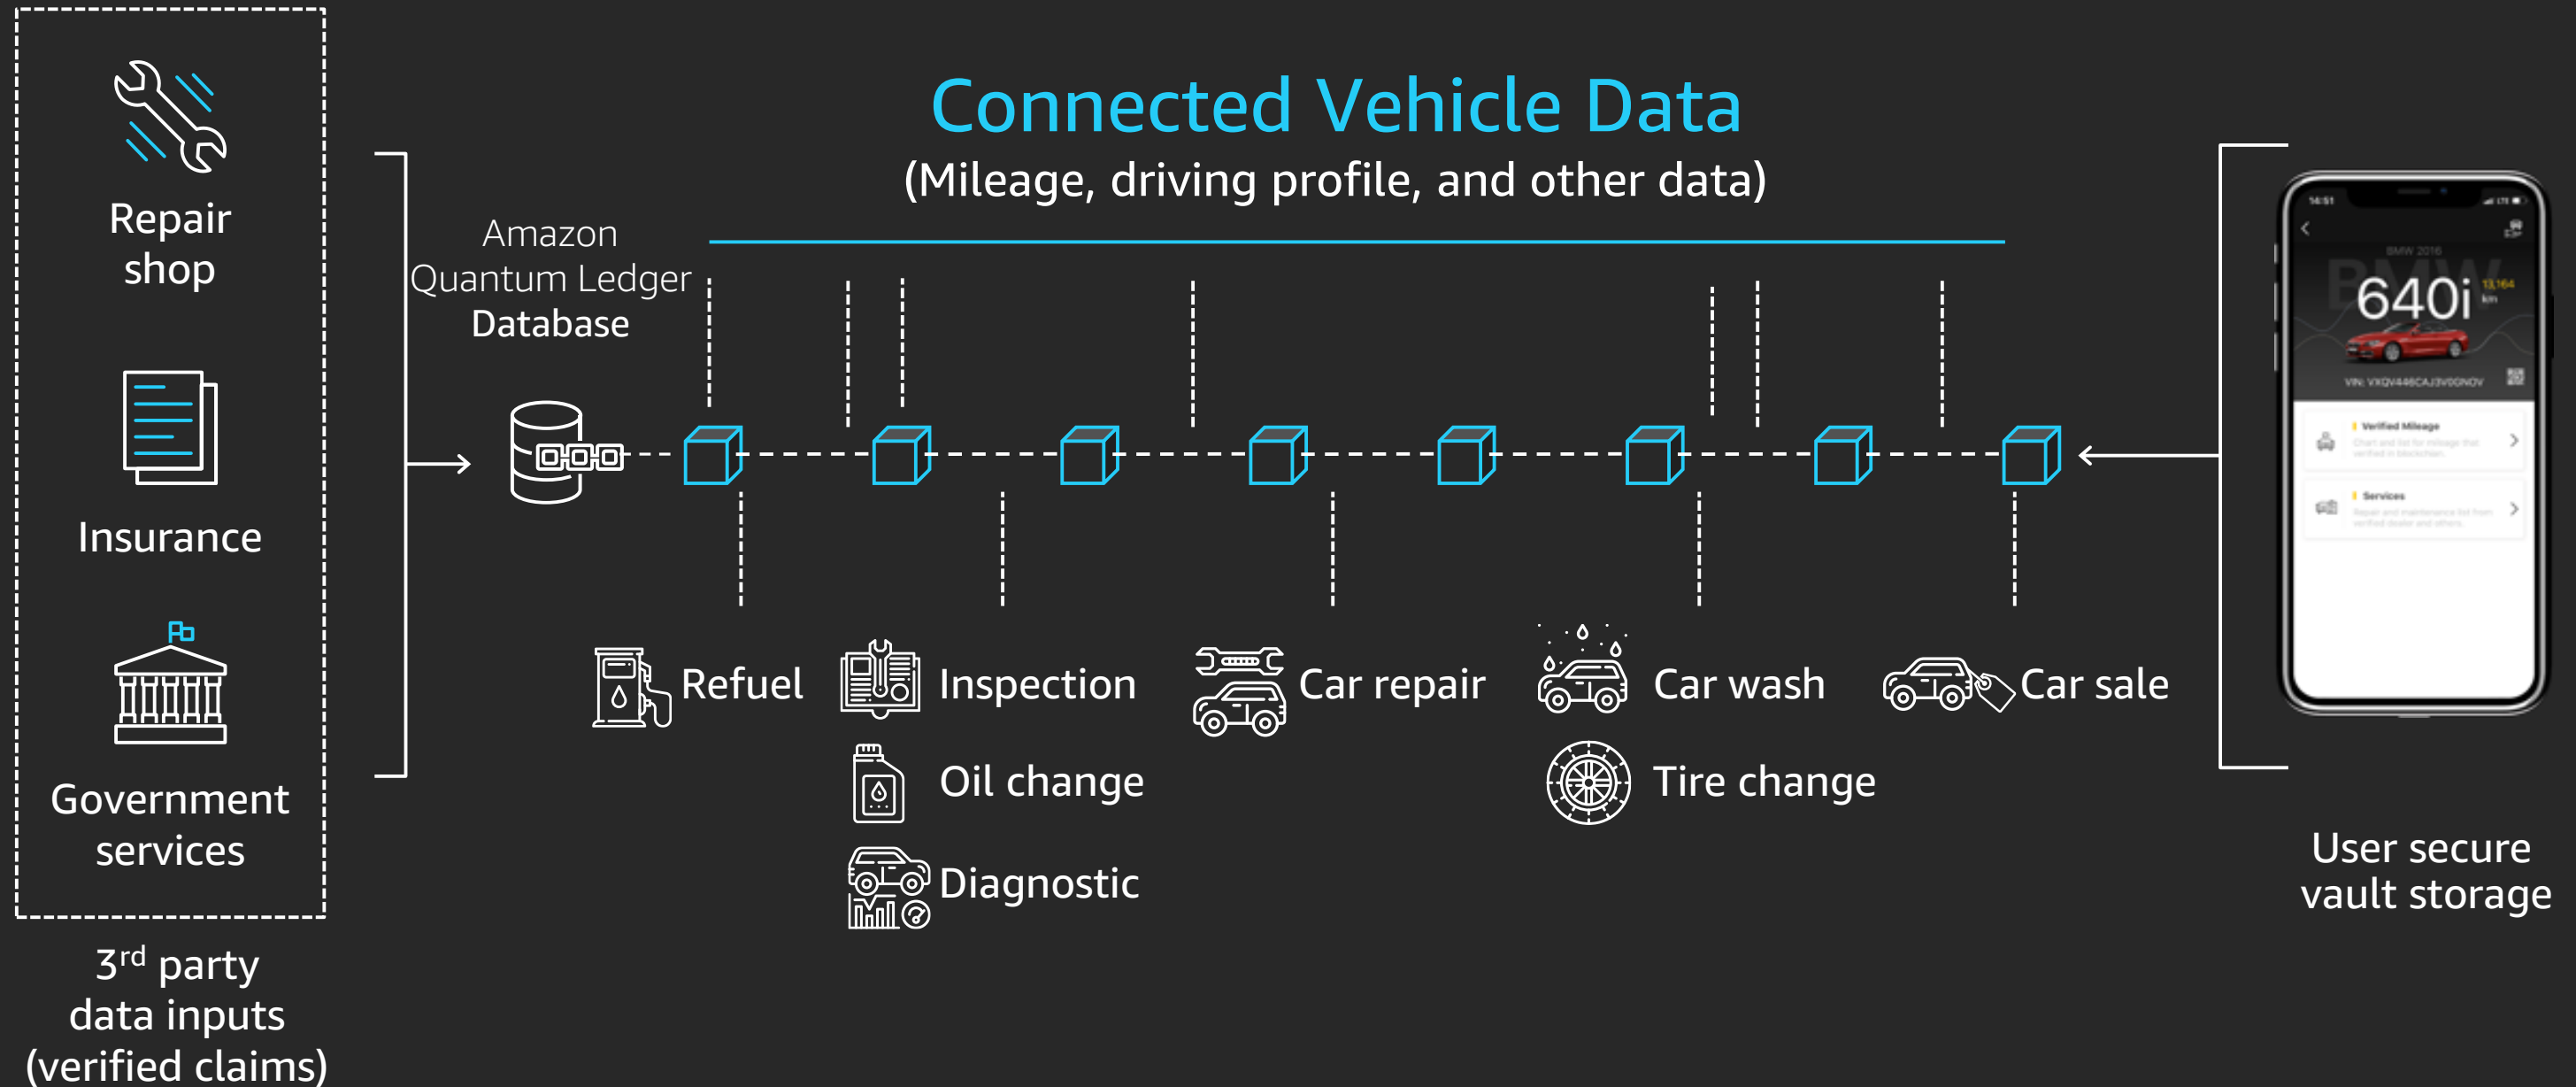


Ensure customer
privacy

Cryptographic Ledger Technology enables end-to-end data pipelines with
verification built in at every stage

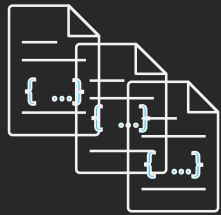
Trusted vehicle data: The foundation of ecosystems

Digital Vehicle Passport Prototype



QLDB as enabler for automotive data ecosystems

Historical data



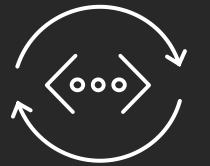
Trusted, verifiable data



On-demand access



Ease of use



Amazon
Quantum
Ledger
Database



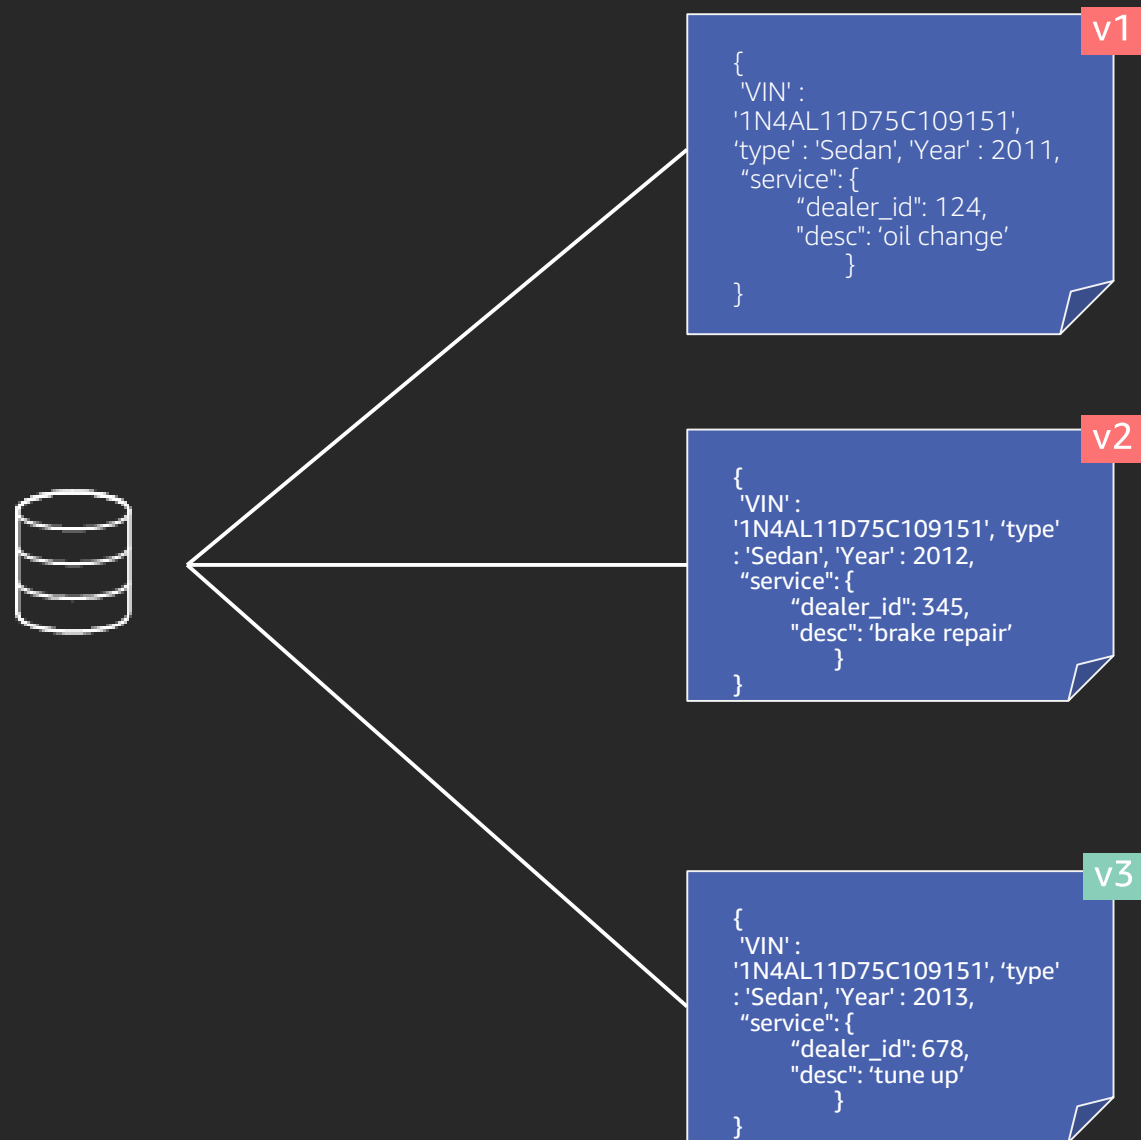
Journaled data revisions

Cryptographic hashing /
verification

Higher throughput
than decentralized
ledgers

Schema-less design:
SQL-like, ACID

The history / versioning problem



Enables:

- Data integrity protection
- Lock access to old document versions
- Disable deletion capability even for admin
- Manually keep indexes / tracking of doc versioning
- Allow external data verification

QLDB provides elegant solution to this problem out-of-the-box

Data model

```
{
  "document_id": "3Lb4dnu5c2Y3Uc04OM",
  "vid": "WBA123456789012345",
  "ownerID": "bmw.driver@bmw.com",

  "vehicleSpecification": {
    "brand": "BMW",
    "model": "I8",
    "productionDate": 2016-08-23,
    "color": "FLUID BLACK MIT AKZENT",
    "driveTrain": "BEV",
    ...
  },

  "lastState": {
    "lastStateCallTime": 1570782292689,
    "geoPosition": "Fi6bnRHN10mCwcR",
    "mileage": "AzLG66DISvu2fGhssN"
  },

  "services": [
    "L2889EwdMA4jtnmR",
    "Tses284nfnwi246k",
    ...
  ],
  ...
}
```

Doc_id: AzLG66DISvu2f5L3lGhssN
Rev_id: 1

Mileage doc

{
 "mileage": {
 "value": 2392,
 "time": "2019-04-19T09:00:55.223Z"
 }
}

Doc_id: AzLG66DISvu2f5L3lGhssN
Rev_id: 2

{
 "mileage": {
 "value": 3598,
 "time": "2019-04-30T09:00:55.223Z"
 }
}

Doc_id: AzLG66DISvu2f5L3lGhssN
Rev_id: 3

{
 "mileage": {
 "value": 4844,
 "time": "2019-05-12T09:00:55.223Z"
 }
}

Doc_id: L2889EwdMA4jtnmR
Rev_id: 1

Services docs

Doc_id: L2889EwdMA4jtnmR
Rev_id: 2

{
 "service": {
 "date": "2019-05-12T09:00:55.223Z",
 "dealer_id": 5790,
 "service_dsc": "tyre rotaion"
 }
}

Doc_id: L2889EwdMA4jtnmR
Rev_id: 3

{
 "service": {
 "date": "2019-05-12T09:00:55.223Z",
 "dealer_id": 124,
 "service_dsc": "oil change"
 }
}

Doc_id: Tses284nfnwi246k
Rev_id: 1

{
 "service": {
 "date": "2019-05-12T09:00:55.223Z",
 "dealer_id": 5790,
 "service_dsc": "tyre rotaion"
 }
}

Any doc (i.e., any new modules added in future)

Doc id: -----
R Doc id: -----
- R Doc id: -----
{ - R Doc id: -----
 - Rev_id: ...
 {

 {...}

Data model with customer encryption

```
{
  "document_id": "3Lb4dnu5c2Y3Uc04OM",
  "vid": "WBA123456789012345",
  "ownerID": "bmw.driver@bmw.com",

  "vehicleSpecification": {
    "brand": "BMW",
    "model": "I8",
    "productionDate": 2016-08-23,
    "color": "FLUID BLACK MIT AKZENT",
    "driveTrain": "BEV",
    ...
  },

  "lastState": {
    "lastStateCallTime": 1570782292689,
    "geoPosition": "Fi6bnRHN10mCwcR",
    "mileage": "AzLG66DISvu2fGhssN"
  },

  "services": [
    "L2889EwdMA4jtnmR",
    "Tses284nfnwi246k",
    ...
  ],
  ...
}
```

```
Doc_id: AzLG66DISvu2f5L3lGhssN
Rev_id: 1
```

Doc_id: AzLG66DISvu2f5L3lGhssN
Rev_id: 2

Mileage doc

```
Doc_id: AzLG66DISvu2f5L3lGhssN
Rev_id: 3
```

QJARYAMIGfMA0GCSqGSib3DQEBAQUAA4GNAD
CBIQKBG/fUvcQJARYAMIGfMA0GCSqGSib3D
QEBAQUAA4GN

```
Doc_id: L2889EwdMA4jtnmR
Rev_id: 1
```

Doc_id: L2889EwdMA4jtnmR
Rev_id: 2

Doc_id: L2889EwdMA4jtnmR
Rev_id: 3

TEPMA0GCSqGS Ib3DQEJARYAM
IGfMA0GCSqGS Ib3DQEBAAQUAA
4GNADCBiQKBgQ/CJ9WRa

Services docs

Doc_id: Tses284nfnwi246k
Rev_id: 1

0GCSq0GCSqGS Ib3DQEBAQUAA4
GNADCBiQK3DQEJARYAMIGfMA0
GCSqGS Ib3DQBgQ CJ9WRanGGS I
b3DQEBAQUAA4GNADCBiQKBgQC
J9WRanG

Any doc (i.e., any new modules added in future)

```
Doc_id: -----
R Doc_id: -----
- R Doc_id: -----
{ - R Doc_id: -----
  - Rev_id: ...
  { -----
    {...}
```

Verifying a record

1.) Capture Docid+Blockaddress of the document

```
Doc_id: L2889EwdMA4jtnmR  
Rev_id: 3
```

```
TEPMA0GCSqGSib3DQEJARYAM  
IGfMA0GCSqGSib3DQEBAQUAA  
4GNADCBiQKBgQ/CJ9WRa
```

```
SELECT r.metadata.id, r.blockAddress FROM  
_ql_committed_VehicleServices AS r  
WHERE ...
```

2.) Retrieve the QLDB Digest



Digest

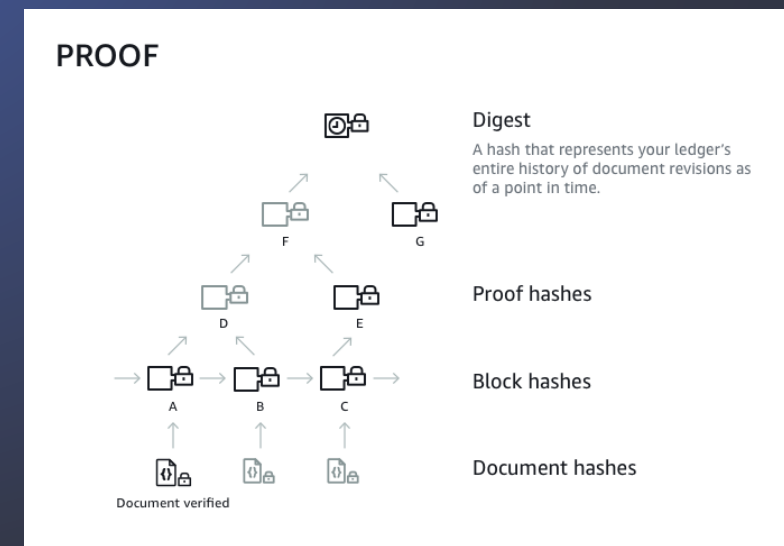
```
{ "digest": "42zaJOfV8iGuthD5Xb/5B9lScHnvxPXm9E=",  
  "digestTipAddress": "{strandId:"BlFTjlSXzOszcE3",sequenceNo:73}",  
  "ledger": "my-ledger",  
  "date": "2019-04-17T16:57:26.749Z" }
```

3.) Execute the verification process (proof)

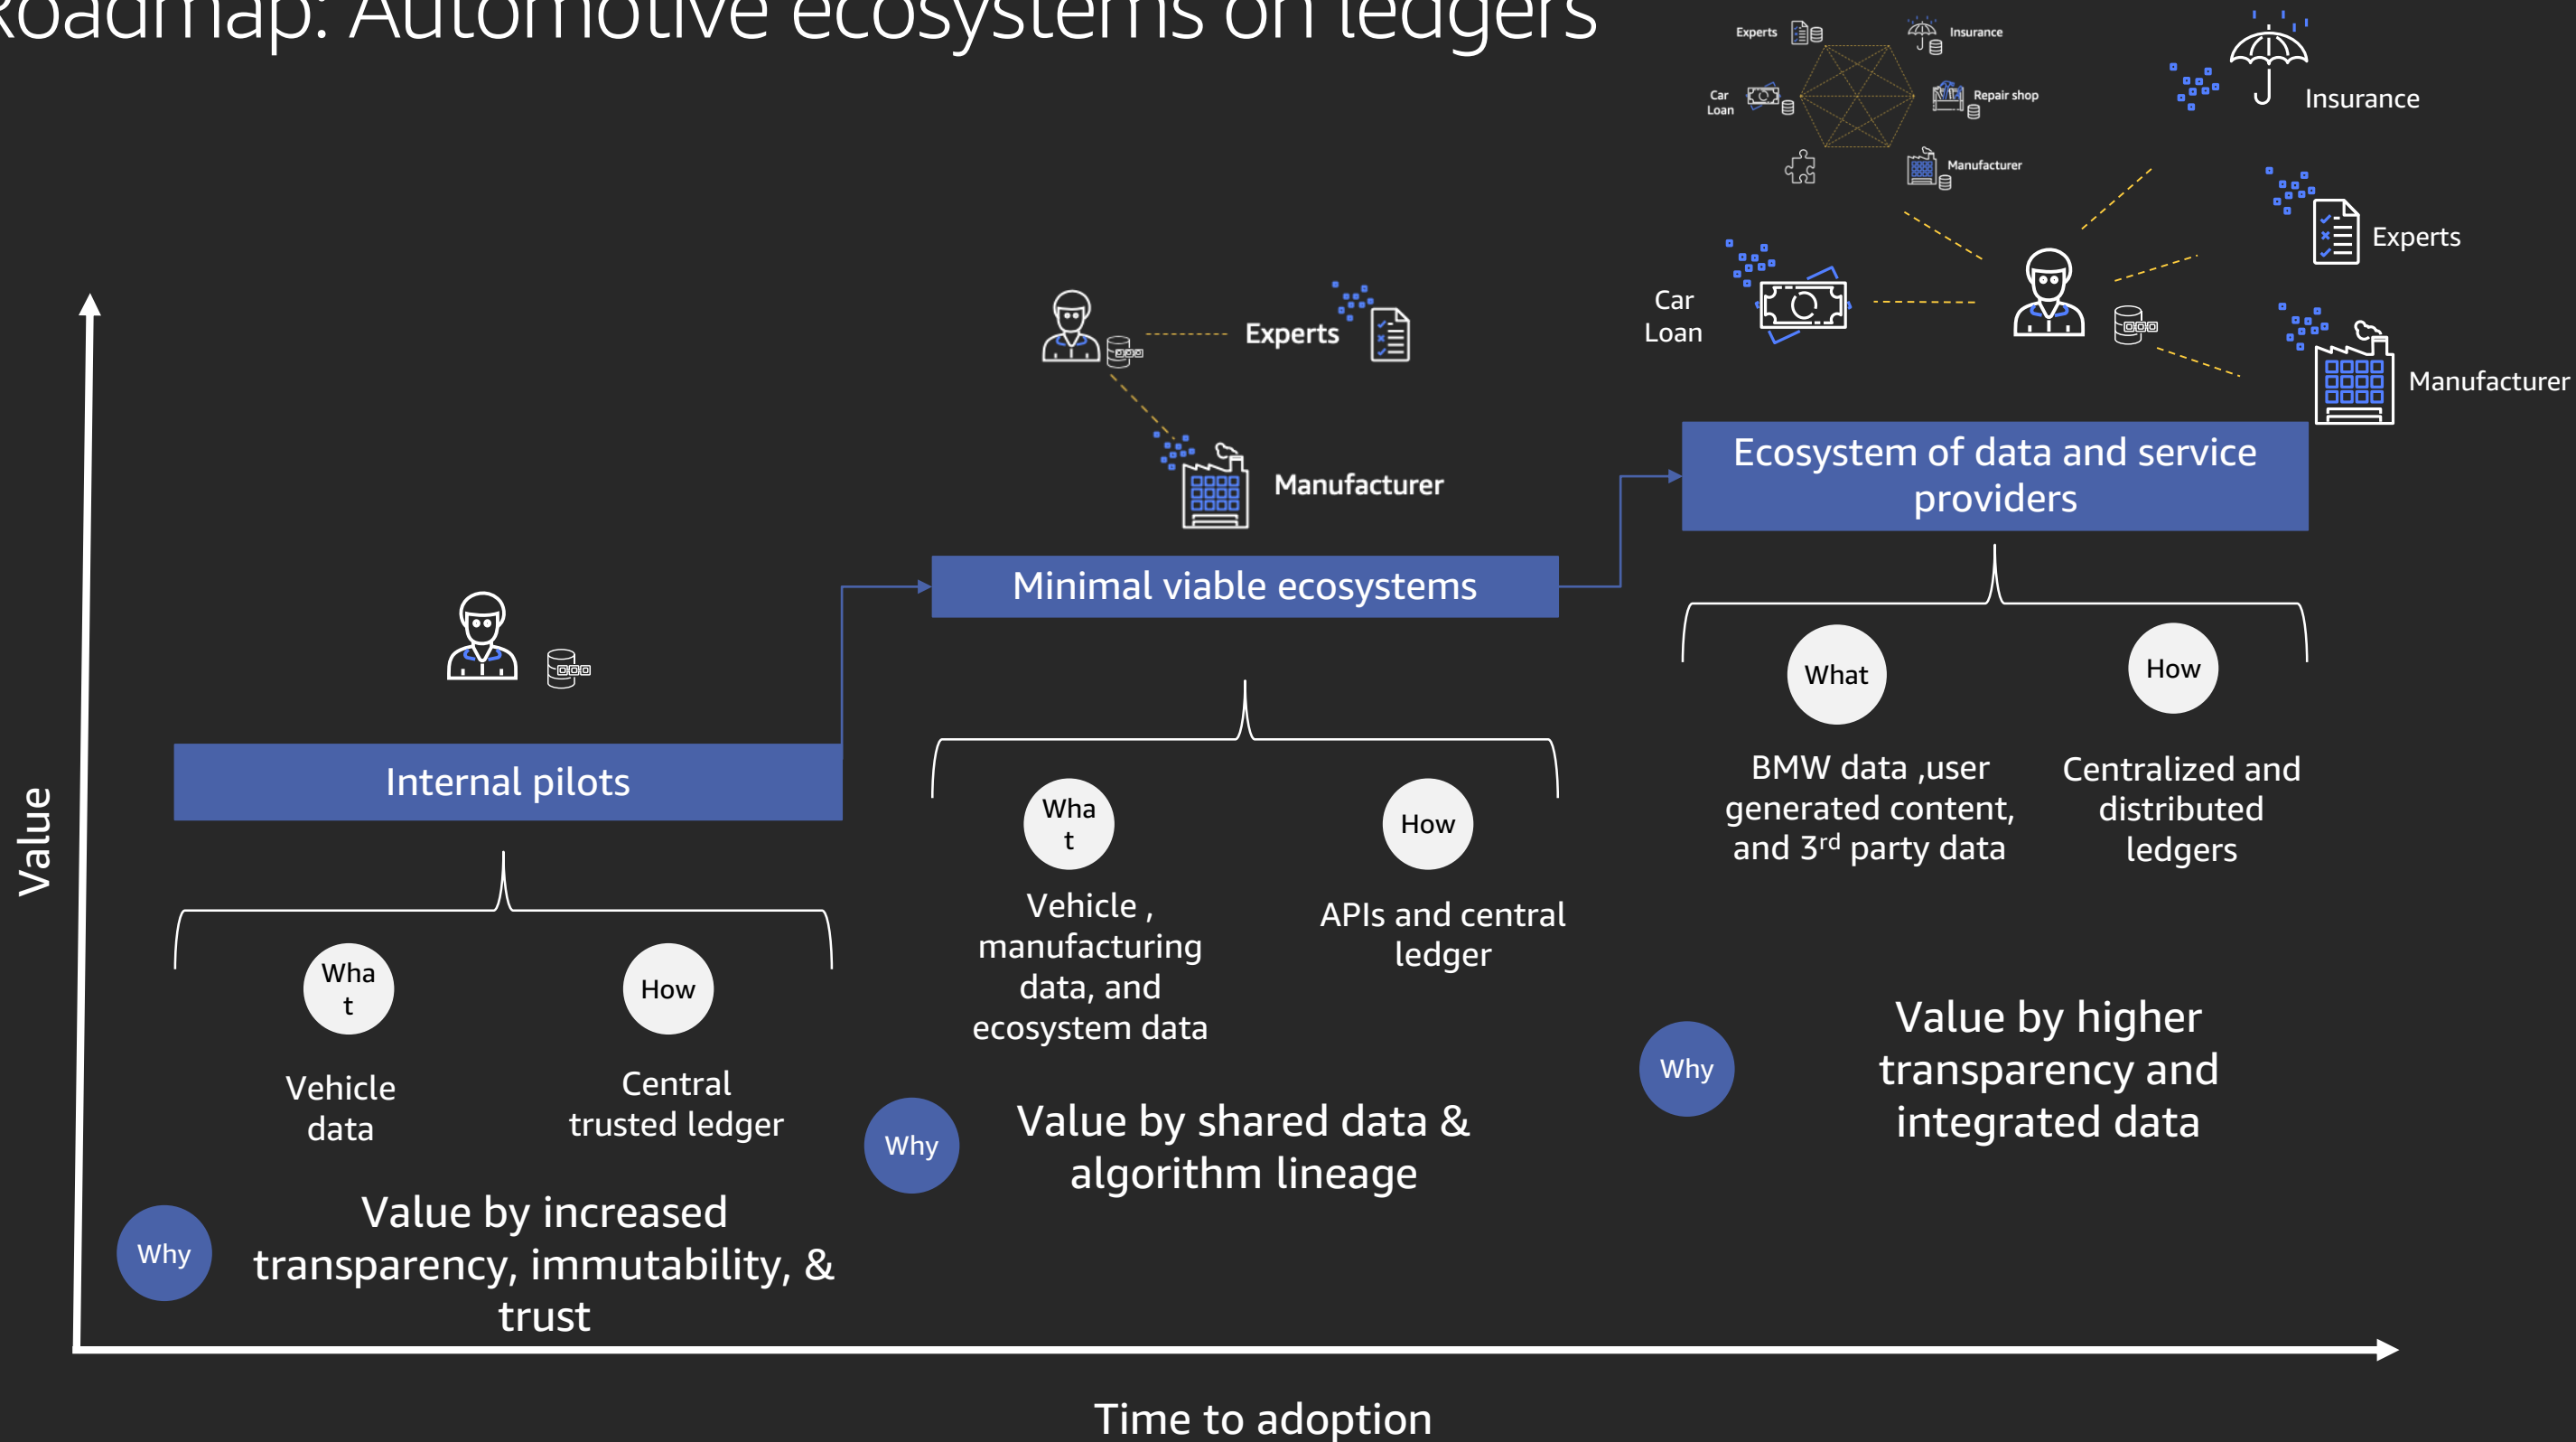
a4e31e36910d99b
d19b7f875f0a0459
7dc0ff52c2f164a16
a9288aed9e710fdd

d07fc3d67314905
dd065d559887900
70410e87072f27ce
2d1cb56457c0879
bc8

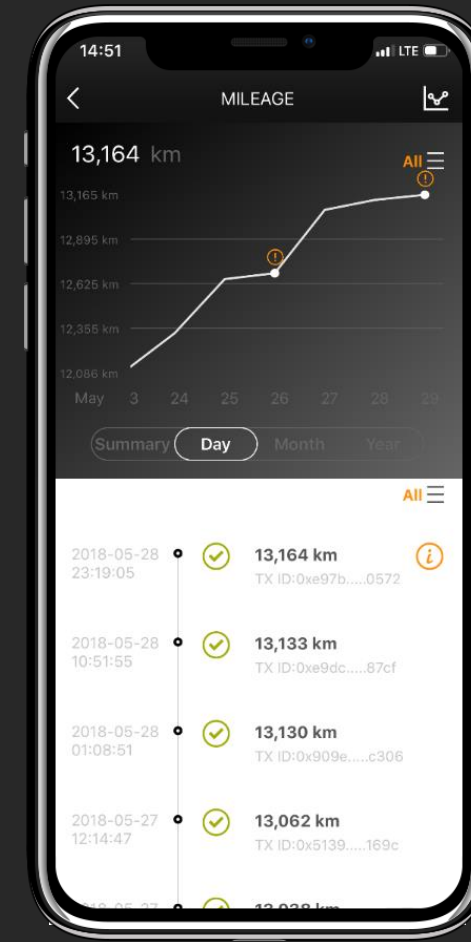
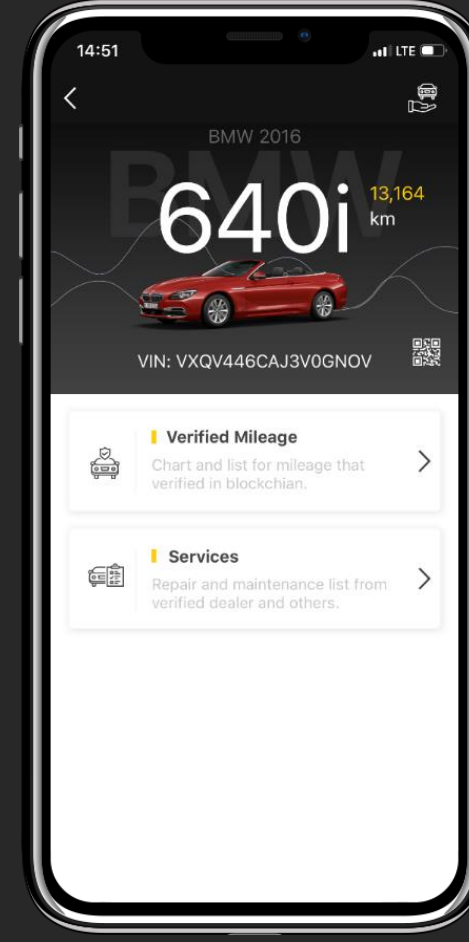
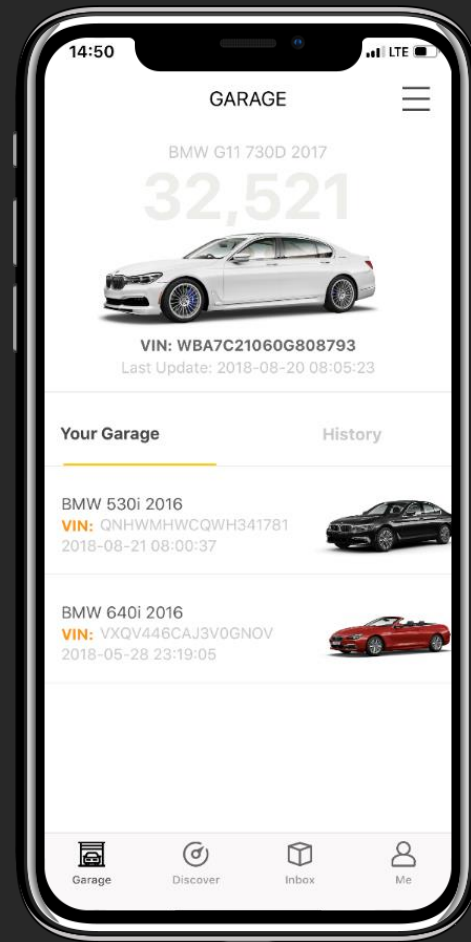
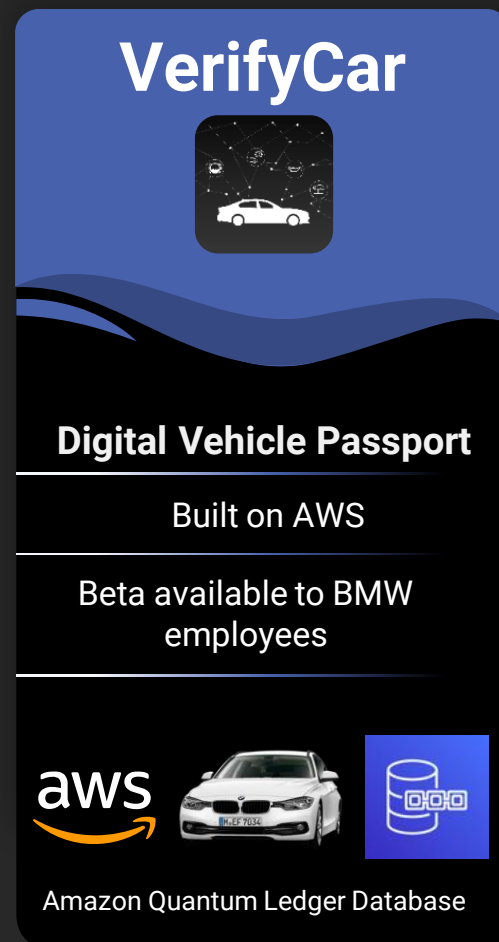
2fc7e994c884bd13
d5fd22b7425328d
0e5d5b0cdcba4d28
5b198be612f229cc
b



Roadmap: Automotive ecosystems on ledgers



Digital vehicle passport application



Thank you!



Please complete the session
survey in the mobile app.